



Delivering Data Science
In Resources & Energy



Knowledge Discovery from Natural Language Processing & Text Mining

DAY 12: Special Data Types

15-Day Data Science Springboard

Assoc Prof Wei Liu & Tyler Bikaun, Department of
Computer Science & Software Engineering, UWA
wei.liu@uwa.edu.au
tyler.bikaun@research.uwa.edu.au

Program partners





Program Timeline

DAY 12: Special Data Types – Natural Language Processing & Text Mining

2-hour Leading Data Scientists Leader Support	Preparatory		Introduction to Data Projects	Data Analysis			Data & Communication Sandbox	Data Fusion and Machine Learning		Data Fusion Sandbox	Special Data Types - Time-series & Networks	Special Data Types - Knowledge Discovery from Natural Language Processing & Text Mining	Special Data	Capstone Project Development & Presentation	Capstone Propeller
2-hours	Day 1	Day 2	Day 3	Day 4	Day 5	Day 6	Day 7	Day 8	Day 9	Day 10	Day 11	Day 12	Day 13	Day 14	Day 15
Enabling your people's data science upskilling & project delivery in 15 day program	Introduction to the program tools	Introduction to the program tools	Zero to Data Science in a day	Getting to know the Program Tools: Data munging and exploratory data analysis	Simple predictions: Regression and statistical model building	Multivariate analysis and model building	Effective data storytelling: Communicating results to non-technical audiences	Pros and cons of commonly used statistical and machine learning techniques I	Pros and cons of commonly used statistical and machine learning techniques II	Consolidate approaches covered and test on datasets	The 4th dimension and predictions	Finding needles in wordstacks	Spatial analytics and predictions	Capstone Project pitches to leadership	Project Review Day



Schedule

DAY 12

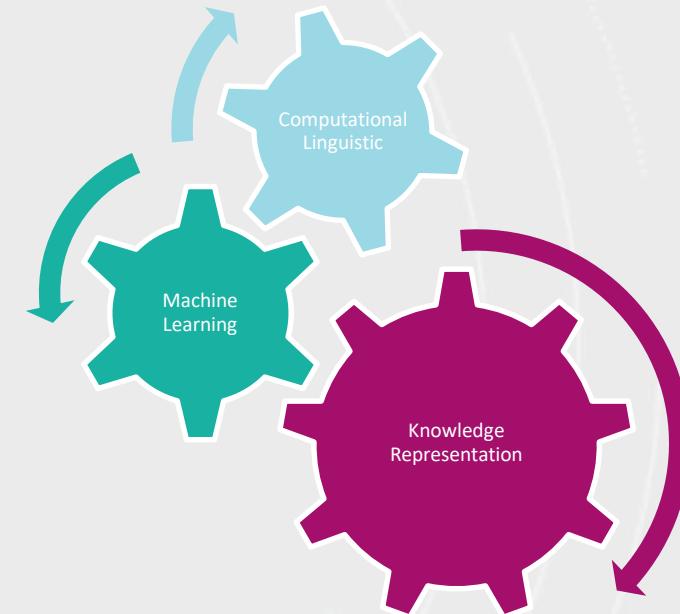
AWST	AEST	Agenda	
08:00	10:00	Open JupyterHub, Q&A	
08:15	10:15	Part 1 (<i>5 min break ~8:45/10:45</i>)	Tyler
<i>09:30</i>	<i>11:30</i>	<i>Morning Tea</i>	
09:45	11:45	Part 2 (<i>5 min break ~10:30/12:30</i>)	Tyler
<i>11:15</i>	<i>13:15</i>	<i>Lunch</i>	
12:00	14:00	Part 3 (<i>5 min break ~12:45/14:45</i>)	Tyler
<i>13:30</i>	<i>15:30</i>	<i>Afternoon Tea</i>	
13:45	15:45	Project Update	Tamryn
14:00	16:00	Part 4 (<i>5 min break ~14:30/16:30</i>)	Tyler
15:15	17:15	Q&A, Reflections, Takeaways, Menti Feedback	Tamryn
15:30	17:30	Close	



A day on Natural Language Processing and Text Mining

What is today about?

- Natural Language Processing
- Language Representation
- Supervised Learning
- Natural Language Annotation
- Unsupervised Learning





Aims & Learning Outcomes – Day 12

Aims

1. Understand practical strategies and applications for using text
2. Provide a foundation for NLP fundamentals including text wrangling, pre-processing, and representation
3. Gain familiarity with supervised and unsupervised learning
4. Understand the natural language annotation process

Learning Outcomes

1. Understand how to wrangle, pre-process and gain insight into text
2. Perform a range of supervised learning tasks
3. Understand the text annotation process
4. Perform unsupervised learning

12.0 Overview

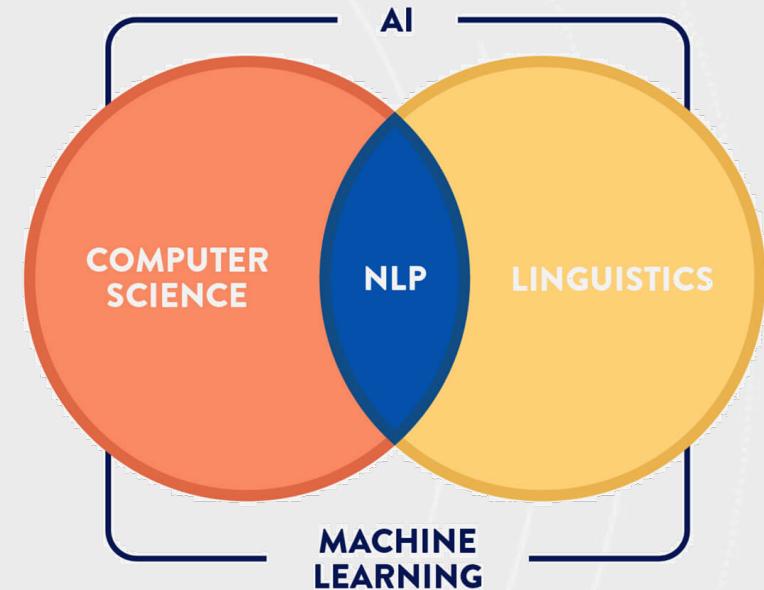
What can we get out of natural language texts



Tasks in Natural Language Processing¹

NLP Tasks

- Data-to-Text Generation
- Grammatical error correction
- Lexical normalization
- Machine translation
- Named entity recognition
- Question answering
- Relation extraction
- Sentiment Analysis
- Simplification
- Summarization
- Text classification
- ...

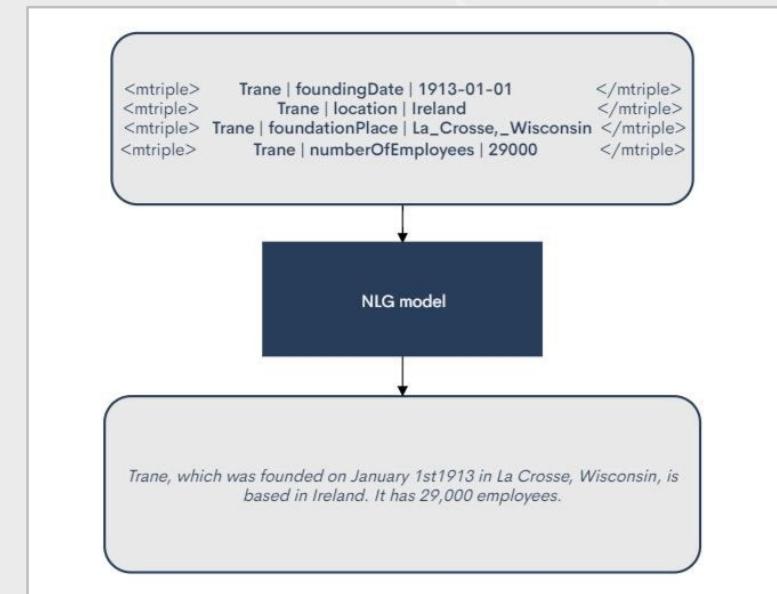




Tasks in Natural Language Processing¹

NLP Tasks

- **Data-to-Text Generation**
- Grammatical error correction
- Lexical normalization
- Machine translation
- Named entity recognition
- Question answering
- Relation extraction
- Sentiment Analysis
- Simplification
- Summarization
- Text classification
- ...





Tasks in Natural Language Processing¹

NLP Tasks

- Data-to-Text Generation
- **Grammatical error correction**
- Lexical normalization
- Machine translation
- Named entity recognition
- Question answering
- Relation extraction
- Sentiment Analysis
- Simplification
- Summarization
- Text classification
- ...

Input (Erroneous)	Output (Corrected)
A important part of my life have been a people that stood by me.	An important part of my life has been the people who stood by me.



Tasks in Natural Language Processing¹

NLP Tasks

- Data-to-Text Generation
- Grammatical error correction
- **Lexical normalization**
- Machine translation
- Named entity recognition
- Question answering
- Relation extraction
- Sentiment Analysis
- Simplification
- Summarization
- Text classification
- ...

source: got **exo** to share, **u** interested? Concert in **hk** !

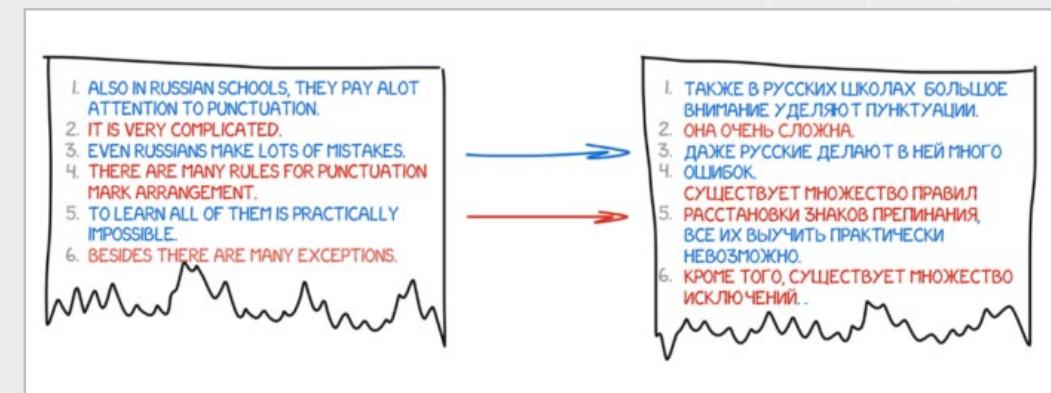
target: got **extra** to share, **are you** interested? Concert in **hong kong** !



Tasks in Natural Language Processing¹

NLP Tasks

- Data-to-Text Generation
- Grammatical error correction
- Lexical normalization
- **Machine translation**
- Named entity recognition
- Question answering
- Relation extraction
- Sentiment Analysis
- Simplification
- Summarization
- Text classification
- ...





Tasks in Natural Language Processing¹

NLP Tasks

- Data-to-Text Generation
- Grammatical error correction
- Lexical normalization
- Machine translation
- **Named entity recognition**
- Question answering
- Relation extraction
- Sentiment Analysis
- Simplification
- Summarization
- Text classification
- ...



Tasks in Natural Language Processing¹

NLP Tasks

- Data-to-Text Generation
- Grammatical error correction
- Lexical normalization
- Machine translation
- Named entity recognition
- **Question answering**
- Relation extraction
- Sentiment Analysis
- Simplification
- Summarization
- Text classification
- ...

Passage Sentence

In meteorology, precipitation is any product of the condensation of atmospheric water vapor that falls under gravity.

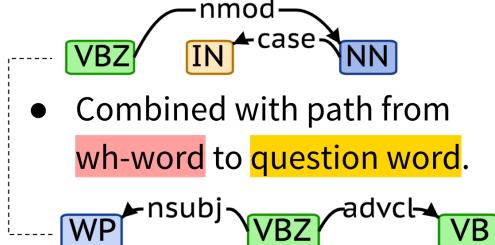
Question

What causes precipitation to fall?

Answer Candidate

gravity

- Path from **passage sentence** words (that also occur in **question**) to **answer**

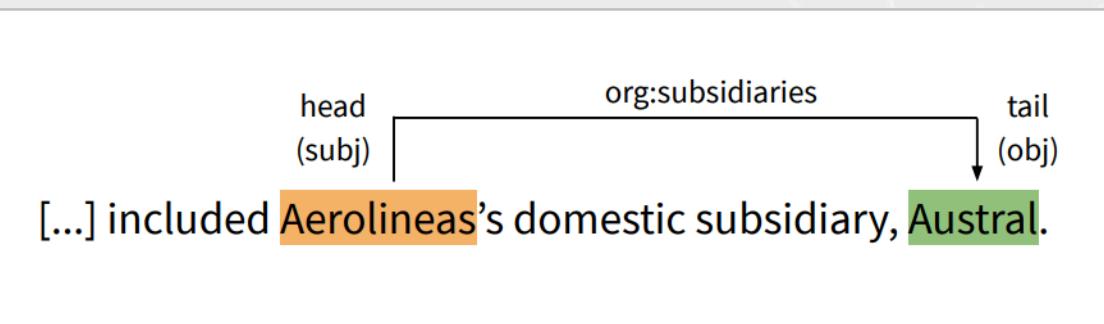




Tasks in Natural Language Processing¹

NLP Tasks

- Data-to-Text Generation
- Grammatical error correction
- Lexical normalization
- Machine translation
- Named entity recognition
- Question answering
- **Relation extraction**
- Sentiment Analysis
- Simplification
- Summarization
- Text classification
- ...





Tasks in Natural Language Processing¹

NLP Tasks

- Data-to-Text Generation
- Grammatical error correction
- Lexical normalization
- Machine translation
- Named entity recognition
- Question answering
- Relation extraction
- **Sentiment Analysis**
- Simplification
- Summarization
- Text classification
- ...



HP Officejet 6500A Plus e-All-in-One Color Ink-jet - Fax / copier / printer / scanner
\$89 online, \$100 nearby ★★★★ 377 reviews
September 2010 - Printer - HP - Inkjet - Office - Copier - Color - Scanner - Fax - 250 sh

Reviews

Summary - Based on 377 reviews

1 star 2 3 4 stars 5 stars

Topic	Rating	Review
ease of use	1 star	"This was very easy to setup to four computers."
value	1 star	"Appreciate good quality at a fair price."
setup	1 star	"Overall pretty easy setup."
customer service	1 star	"I DO like honest tech support people."
size	2 stars	"Pretty Paper weight."
mode	2 stars	"Photos were fair on the high quality mode."
colors	2 stars	"Full color prints came out with great quality."



Tasks in Natural Language Processing¹

NLP Tasks

- Data-to-Text Generation
- Grammatical error correction
- Lexical normalization
- Machine translation
- Named entity recognition
- Question answering
- Relation extraction
- Sentiment Analysis
- **Simplification**
- Summarization
- Text classification
- ...

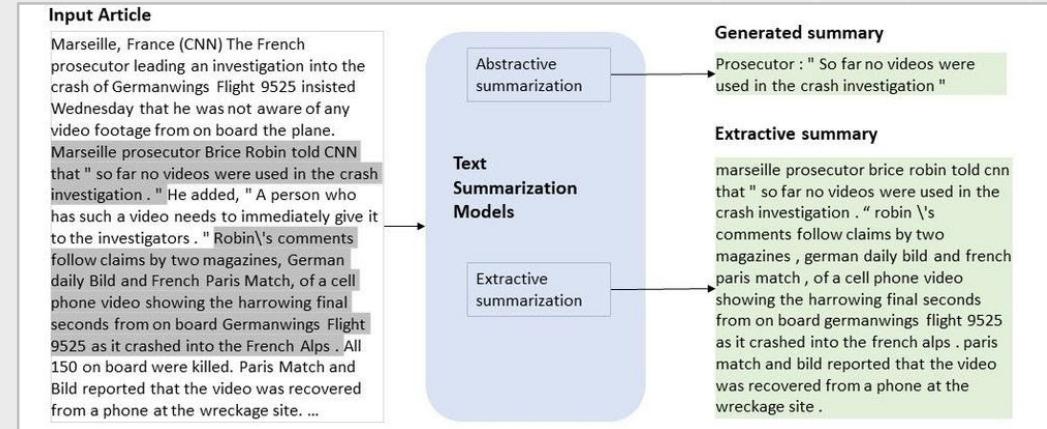
Example of a Complex Sentence	Example of a Simplified Sentence
Grammarly provides assistance in order to optimize users' communication.	Grammarly helps people communicate.



Tasks in Natural Language Processing¹

NLP Tasks

- Data-to-Text Generation
- Grammatical error correction
- Lexical normalization
- Machine translation
- Named entity recognition
- Question answering
- Relation extraction
- Sentiment Analysis
- Simplification
- **Summarization**
- Text classification
- ...

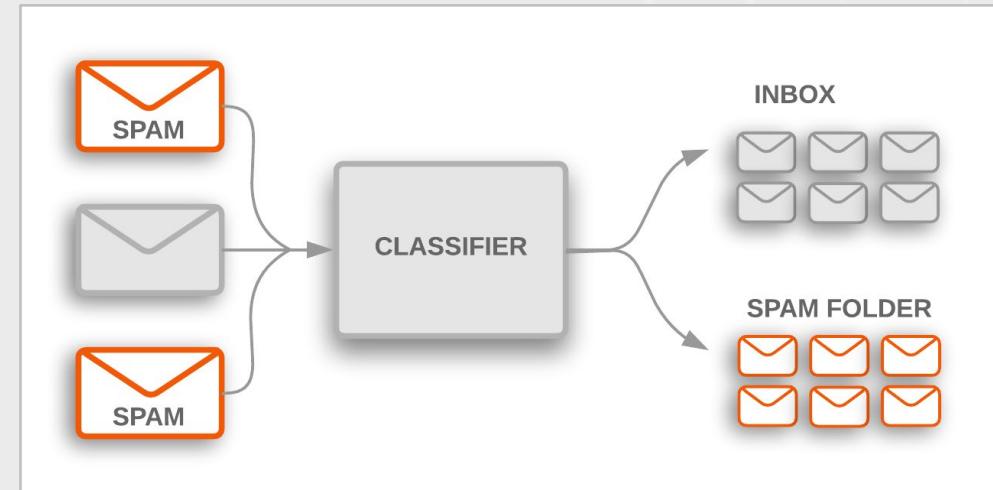




Tasks in Natural Language Processing¹

NLP Tasks

- Data-to-Text Generation
- Grammatical error correction
- Lexical normalization
- Machine translation
- Named entity recognition
- Question answering
- Relation extraction
- Sentiment Analysis
- Simplification
- Summarization
- **Text classification**
- ...



Our Work and Projects



Programming Language or Natural Language?

Natural Language

- General
- Grammatically Correct
- Large Vocabulary
- Not noisy
- Examples: Wikipedia

Technical Language

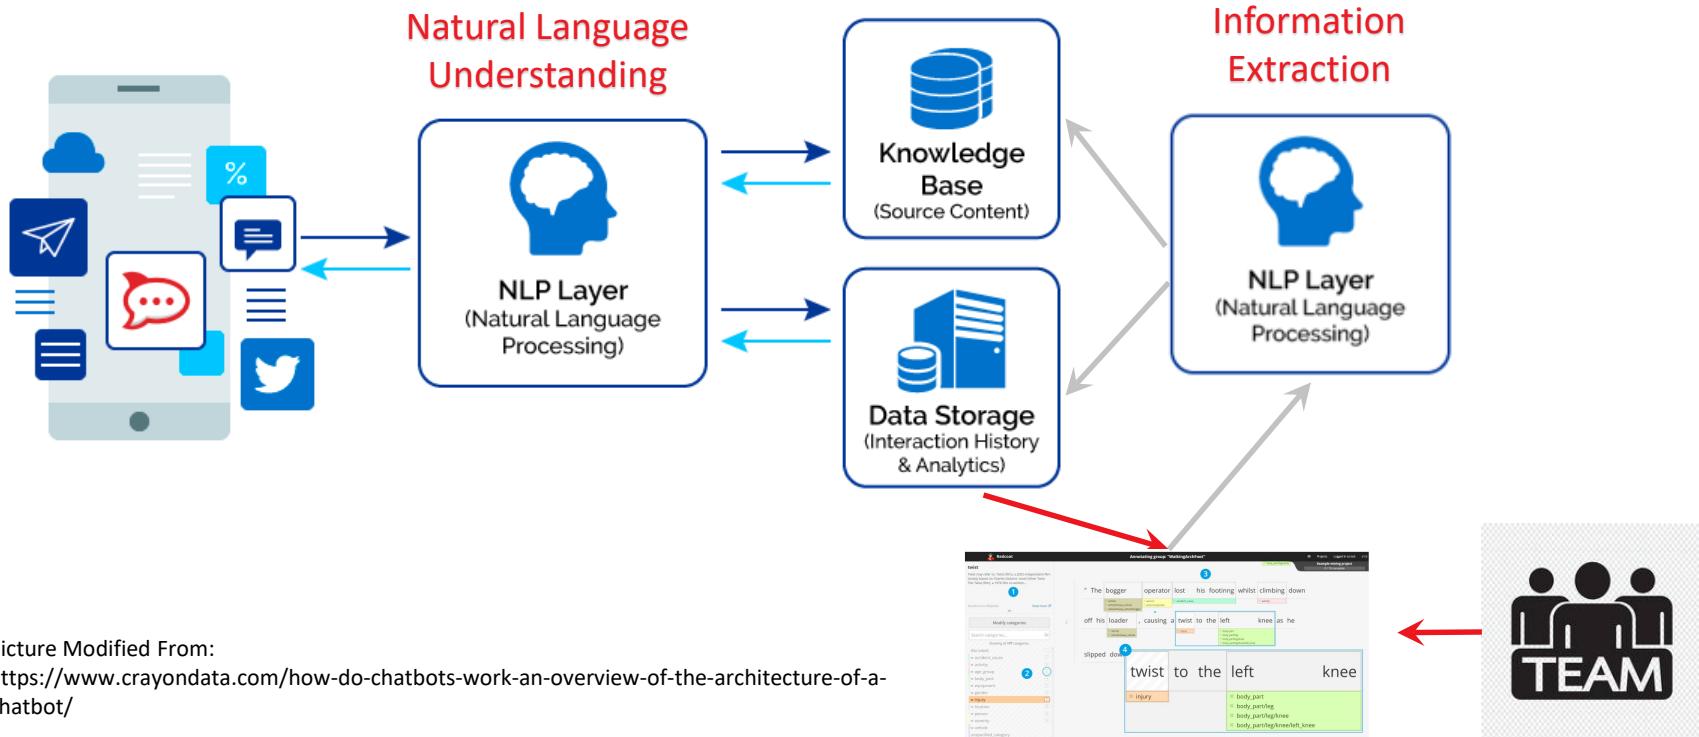
- ✓ Domain Specific
- ✓ Small Vocabulary
- ✗ Terse
- ✗ Jargons, Misspelling
- Examples: Safety logs, Maintenance work orders, geological reports

Technical language processing (TLP) aims to understand **technical languages**, a subset of natural languages that appears in industry-specific contexts.

- Lexical Normalisation
- Grammatical Error Correction
- No pretrained embeddings
- No ground truth
- Low-resource datasets

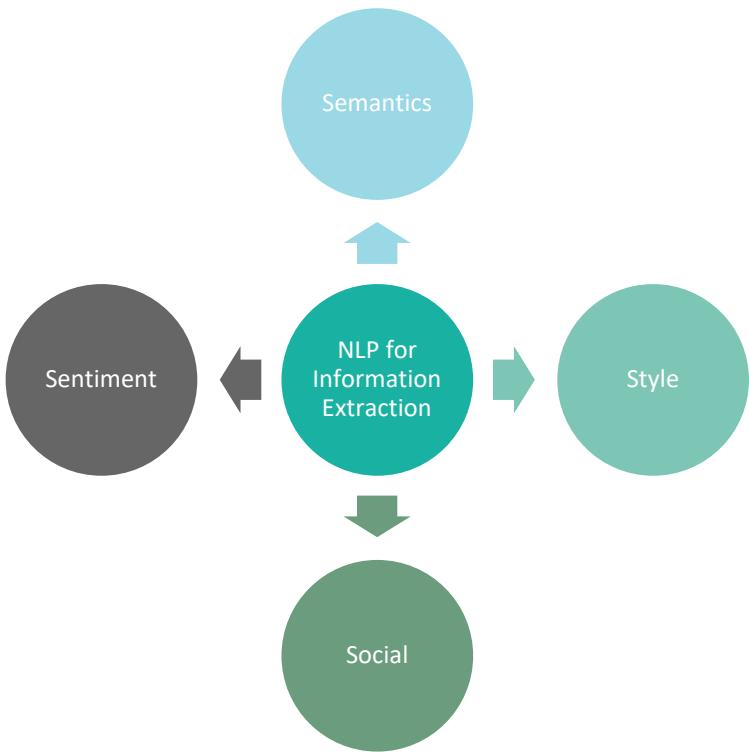


Language – Natural or Technical?





What can we extract out of text? – S4



What, When, Where, How and Why?

Semantics (Entities)

- Entity Recognition

Social (Relations)

- Relation Extraction
- Triple Extraction ([Knowledge Graphs](#))

Style

- Stylometry (fraud, authorship, forensics)
- Grammar (Industry Synthetic Data Generation)

Sentiment

- Social Reviews (e.g. Fraud Detection)



Picture Credit: <https://research.aimultiple.com/chatbot-architecture/>



The hidden treasure in your dataset

W	X	Y	Z	AA	AB	AC	AD	AE	AF	AG	
1	Place Of Injury	Activity	Injury Typ	Body Part	Nature Of	Occupatio	Work Typ	Company	Contractor	Company	Detailed Description
2	0471 - Processing Pla	8230 - Rep	2510 - C/B	6300 - He	5600 - Cru	719000 - E	Full Time	Contracto	Contractor		Traumatic crush injury to the head and chest.
3	0610 - W.Shop Heavy	8130 - Rep	2510 - C/B	6500 - Cor	5400 - Fra	631000 - F	Full Time	Contracto	Contractor		Two fitters were undertaking maintenance work when an accident occurred resulting in
4	0451 - Tailings Stora	7900 - Mis	2130 - S/B	6920 - Mu	5600 - Cru	411000 - P	Full Time	Contracto	Contractor		At approximately 9.20am on Wednesday 4th December a contractor working on the ta
5	0610 - W.Shop Heavy	2110 - Wa	1210 - Fall	5300 - Kne	5200 - Stra	631000 - F	Full Time	Contracto	Contractor		Drill fitter had finished planned services for the shift and had returned the tools to the
6	0370 - Storage Yard/	2450 - Mo	9100 - Bite	5400 - Lov	6310 - Bite	359000 - N	Full Time	Contracto	Contractor 1		The injured person identified a stinging sensation on his left shin area. The injured per
7	0471 - Processing Pla	7900 - Mis	9100 - Bite	5200 - Up	6310 - Bite	359000 - N	Full Time	Contracto	Contractor 1		Injured person was sleeping when he felt a stinging sensation on his right thigh area. T
8	0464 - Reagent / Raw	2450 - Mo	2510 - C/B	4630 - Fin	6100 - Lac	411000 - P	Full Time	Contracto	Contractor 10		An Operator was transporting signs in a loader bucket. Whilst removing one of the sign
9	0464 - Reagent / Raw	2320 - Lifti	4200 - O.X	3100 - Bac	5200 - Stra	372000 - C	Full Time	Contracto	Contractor 10		On 11/11/13 at approximately 0900hrs employee was assisting two other operators in
10	0128 - Level (Develop	2410 - Pull	1120 - Fall	4620 - Har	5200 - Stra	241000 - E	Full Time	Contracto	Contractor 11		Operator was climbing onto the rear of the bogger to straighten a hand rail when it gave
11	0128 - Level (Develop	2180 - Get	3150 - Ste	5500 - Ank	5200 - Stra	269000 - L	Full Time	Contracto	Contractor 11		Service crew member stepped down from work cage onto uneven ground and rolled le
12	0950 - Car Park	2110 - Wa	3120 - Ste	5500 - Ank	5200 - Stra	165000 - S	Full Time	Contracto	Contractor 11		walking in car park from office to workshop & rolled ankle
13	0128 - Level (Develop	2180 - Get	1420 - Fall	5300 - Kne	5200 - Stra	241000 - E	Full Time	Contracto	Contractor 11		The bogger operator lost his footing whilst climbing down off his loader, causing a twi
14	0130 - Underground	2110 - Wa	1232 - Fall	3181 - Bac	6800 - Pail	241000 - E	Full Time	Contracto	Contractor 11		Person was walking around the bogger in the workshop and trod in some grease that v

What, When, Where, How and
Why?

- Individual vs. Population
- Query vs. Machine Learning

activity

walking in car park from office to workshop

location

and rolled ankle

location

location

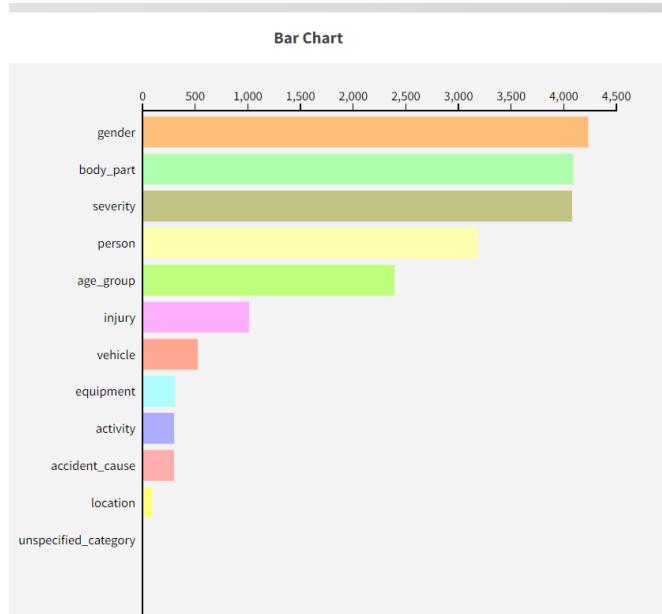
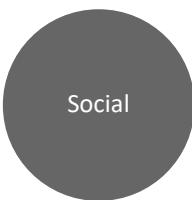
injury

body part





S4 – Semantics and Social

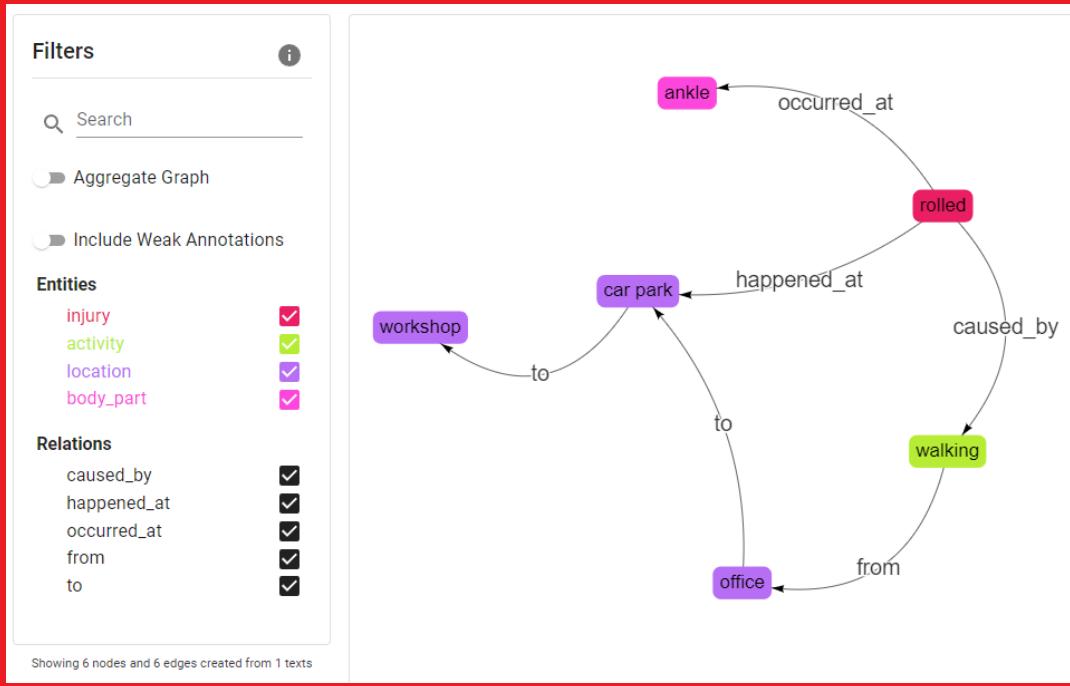


Antecedent	Consequent	Lift	Confidence
operator , serious , haul_truck	➔ neck , back	62.792	0.625
age_15_24 , vehicle	➔ light_vehicle	53.821	1.000
left_eye , minor	➔ male , eye	49.141	0.750
left_eye , minor	➔ eye	45.210	0.750
left_eye , male , minor	➔ eye	45.210	0.750
operator , serious , haul_truck	➔ age_45_54 , neck	39.245	0.625
serious , age_45_54 , haul_truck	➔ operator , neck	38.040	0.833
serious , haul_truck , back	➔ operator , female	34.710	0.714
serious , haul_truck , back	➔ operator , neck	32.606	0.714
operator , serious , haul_truck	➔ back , female	28.542	0.625
operator , age_45_54 , haul_truck	➔ serious , neck	23.685	0.833



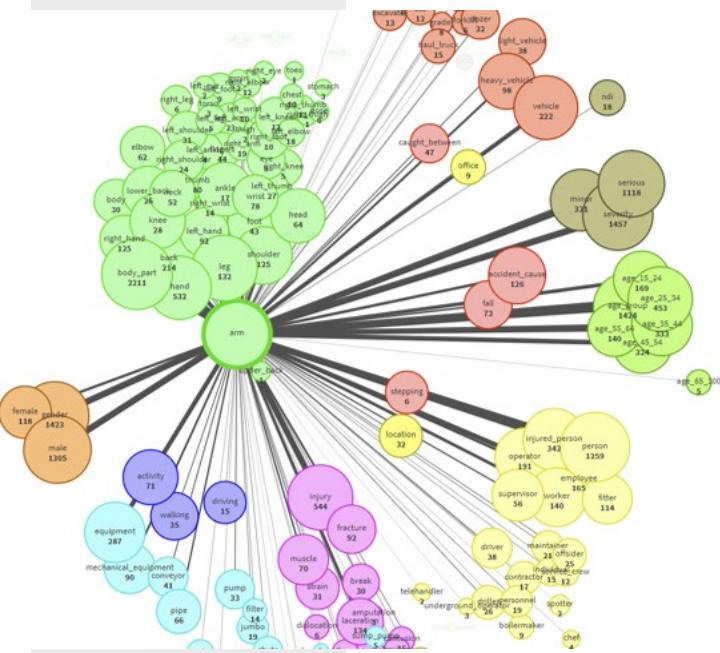
Annotation Tools Developed at UWA – QuickGraph and Lexiclean

Maintenance Work Order Processing - ARC Training Centre for Transforming Maintenance through Data Science

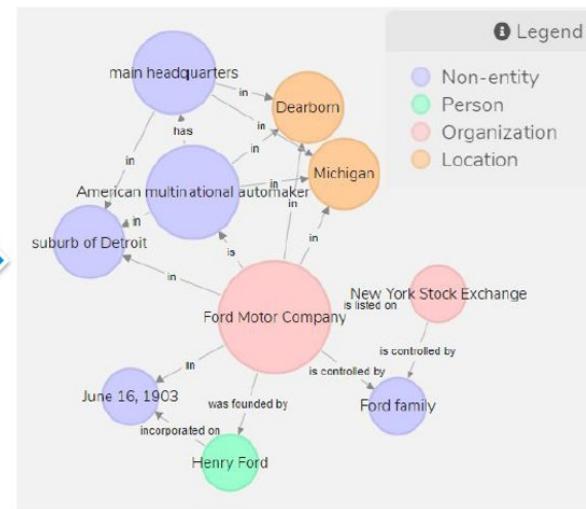




Award Winning NLP & KG Solutions



WAITTA INCITE Award 2019



ICDM Knowledge Graph Contest 1st Prize

UWA Natural Language and Technical Language Processing Group

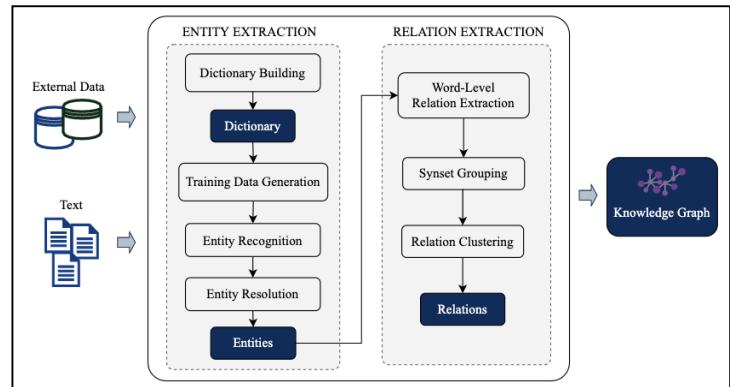
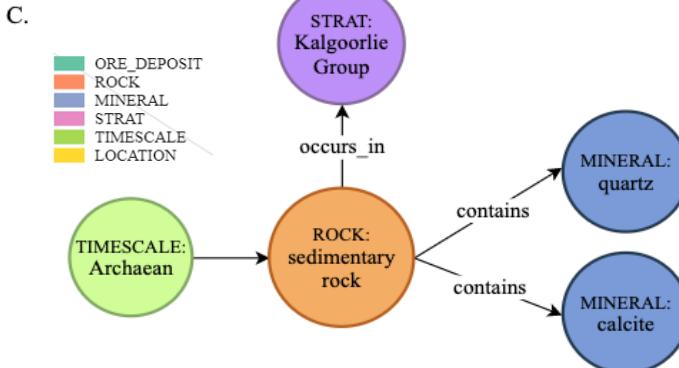
<https://nlp-tlp.org>



KG Extraction from Geological Surveys

- A. Archaean sedimentary rocks occurred within the Kalgoorlie Group.
Most sedimentary rocks contain either quartz or calcite.

- B. Archaean → sedimentary rocks
 sedimentary rocks → occurred within → Kalgoorlie Group
 sedimentary rocks → contain → quartz
 sedimentary rocks → contain → calcite



Knowledge Graph Schema

Enksaikhan et al. (2021), "Understanding Mineral Depositional Conditions using Machine Reading of Text", *Ore Geology Reviews*

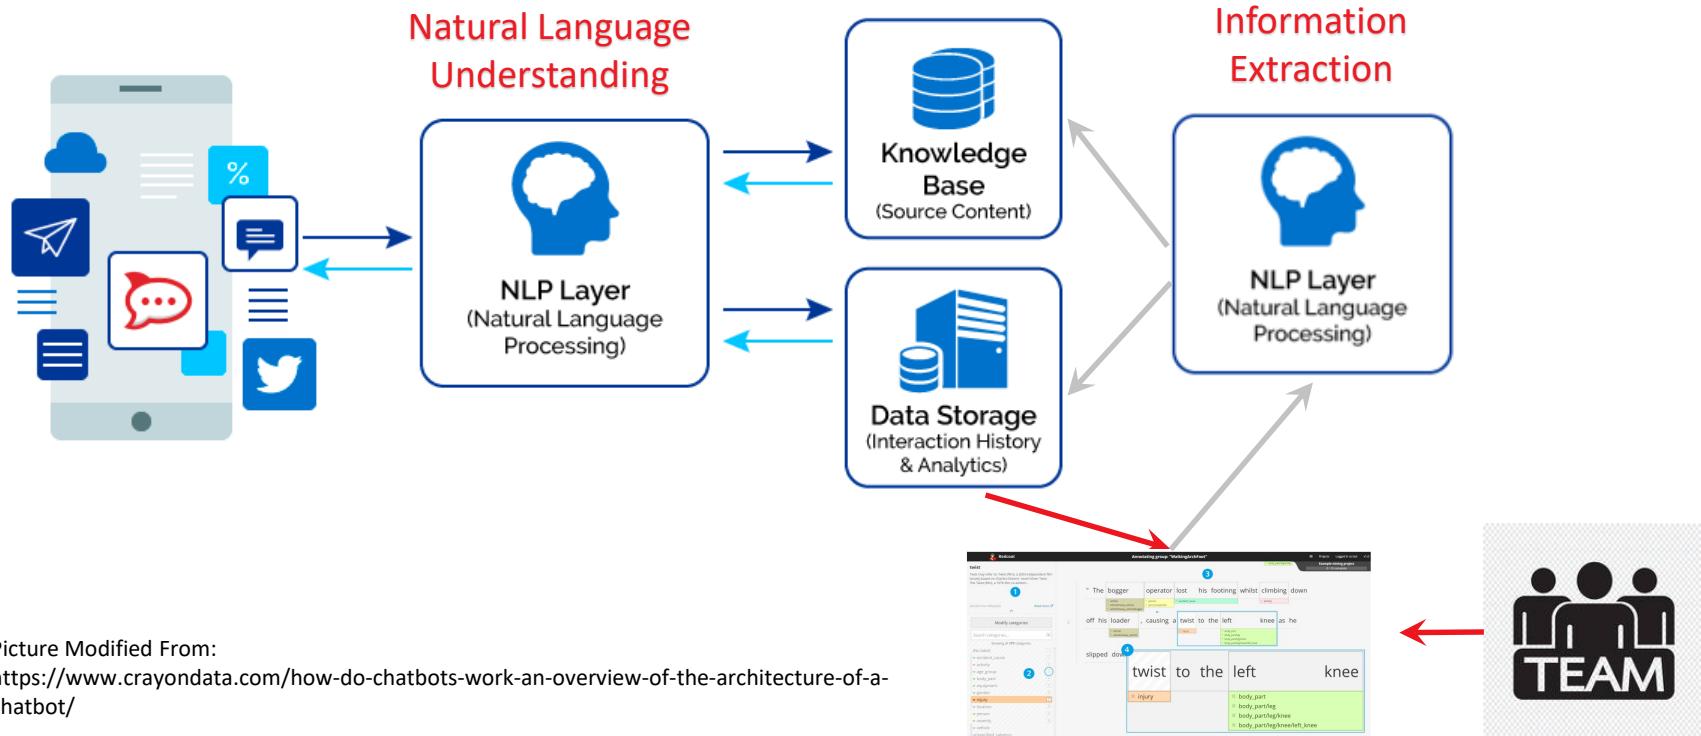
Enksaikhan et al. (2020), "Auto-Labelling Entities in Low Resource Text: A Geological Case Study", *Knowledge and Information Systems*

Enksaikhan et al. (2018), "Towards geological knowledge discovery using vector-based semantic similarity", *Proceedings of the 14th Int. Conf. on Advanced Data Mining and Applications*





TLP – redefining the workforce

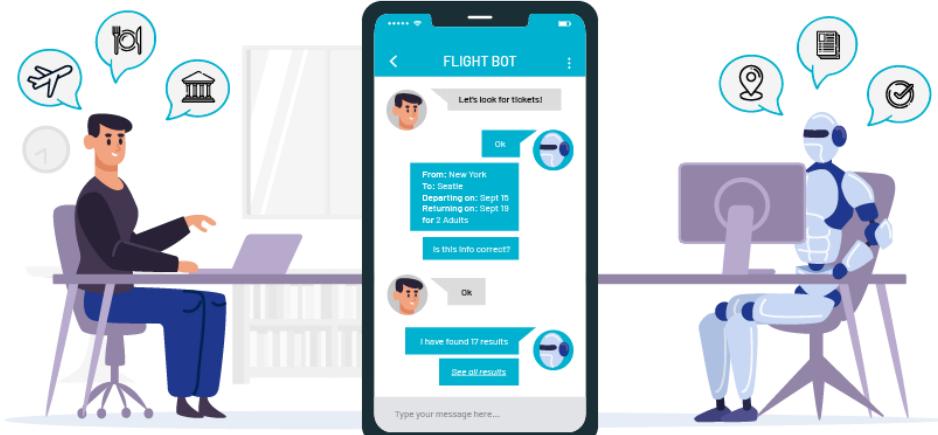


Picture Modified From:

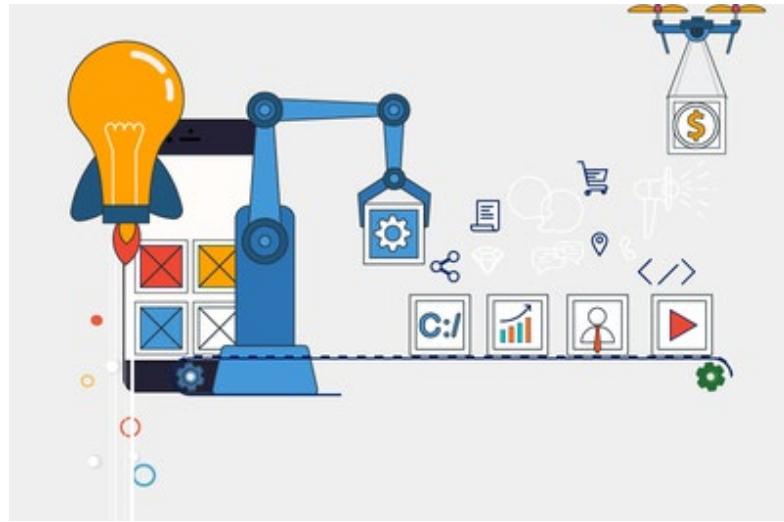
<https://www.crayondata.com/how-do-chatbots-work-an-overview-of-the-architecture-of-a-chatbot/>



Human – Digital Workforce: Trends in Automation 2022



People will work side-by-side with **virtual robotic assistants**—sharing tasks, passing work back and forth, collaborating.



Within the next 5 years, hybrid human-digital workforces will

- Work in a **task-oriented** (rather than software application-oriented) manner, and
- Be part of **virtual assembly line**



THANK YOU



NLP-TLP Group: <https://nlp-tlp.org/>

Academics



Wei Liu

Associate Professor

[UWA Profile](#) | [LinkedIn](#)



Melinda Hodkiewicz

Professor

[UWA Profile](#) | [LinkedIn](#)



Tim French

Senior Lecturer

[UWA Profile](#) | [LinkedIn](#)



Michael Stewart

Postdoctoral Research Fellow

[UWA Profile](#) | [LinkedIn](#)



Eun-Jung Holden

Professor

[UWA Profile](#) | [LinkedIn](#)



Caren Han

Lecturer

[UWA Profile](#)



Naeha Sharif

Lecturer

[UWA Profile](#) | [LinkedIn](#)

PhD Students



Ziyu Zhao

PhD Student

An Efficient Neural Probabilistic Logical Resolution for Multi-class Multi-label Entity Typing

[LinkedIn](#)



Tyler Bikaun

PhD Student

Technical Language Processing for Industrial Maintenance Records

[LinkedIn](#)



Chau Nguyen Duc Minh

PhD Candidate

Query Embedding for Long Reasoning over Natural-Technical Domains



Caitlin Woods

PhD Student

Adaptive User Interfaces for Industrial Maintenance Procedures

[Website](#) | [LinkedIn](#)



Tom Smoker

PhD Student

Rectifying knowledge graph link prediction using embedding-enhanced ontologies

[LinkedIn](#)

AI 2022: <https://ajcai2022.org/>



wei.liu@uwa.edu.au



linkedin.com/in/wei-liu-b0b4521ab

35th Australasian Joint Conference on Artificial Intelligence
5 – 8 December 2022
Perth, Western Australia

12.1 Fundamentals of NLP

Text wrangling, pre-processing and analysis



Fundamentals of Natural Language Processing

Agenda (Hands-on Session) - An introduction to the Natural Language Tool Kit (NLTK)

- Tokenization
- Stop-words
- Feature distributions
- N-grams
- Stemming
- Concordancing
- Dispersion plotting
- Bi-gram significance
- Word contexts
- Word similarities
- Parts-of-Speech (POS)
- Named entity recognition

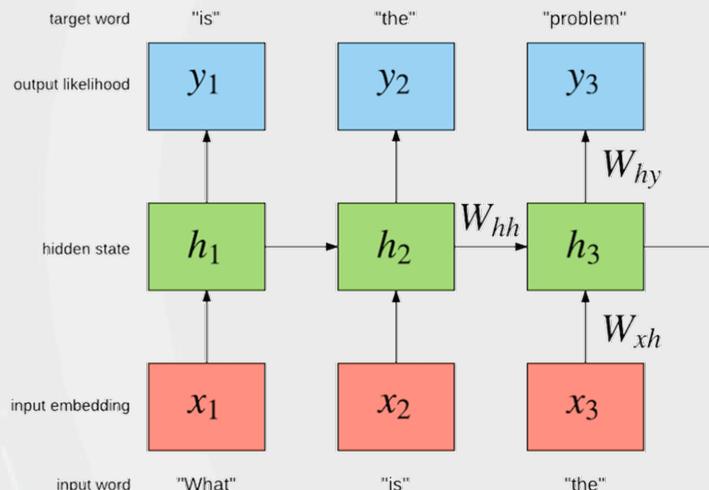
12.2 Fundamentals of NLP

Representing natural language



Word Embeddings: One-hot Vectors

Word representations – one-hot vectors



(Socher, 2018)

What is the name of the prime minister of Australia?

What = [10000000]

is = [01000000]

the = [00100000]

name = [00010000]

of = [00001000]

the = [00100000]

prime = [00000100]

minister = [00000010]

of = [00001000]

Australia = [00000001]



Word Embeddings: Learning Context

Representing words by their context

You shall know a word by the company it keeps

J. R. Firth 1957

...government debt problems turning into **banking** crises as happened in 2009...
...saying that Europe needs unified **banking** regulation to replace the hodgepodge...
...India has just given its **banking** system a shot in the arm...

These context words will represent **banking**

prime = [0,0,0,0,0,1,0,0]
minister = [0,0,0,0,0,0,1,0]



prime = [-0.1, 0.1, 0.2, 0.3, 0.3, 0.4, 0.1, 0.7]
minister = [-0.2, 0.2, 0.4, 0.2, 0.2, 0.2, 0.1, 0.6]

One-hot

Embeddings

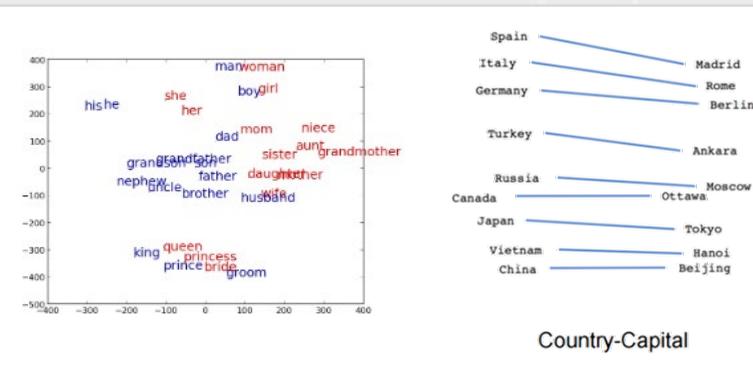
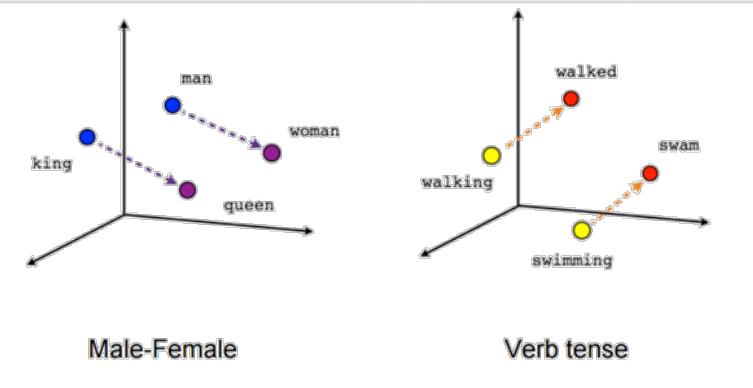


Word Embeddings: Power of Embeddings

Imagine doing simple math with words

“A is to B as C is to D” tasks

$$-\text{word}_A - \text{word}_B - \text{word}_C \cong \text{word}$$





Word Embeddings: Analogy Task Examples

“A is to B as C is to D” tasks

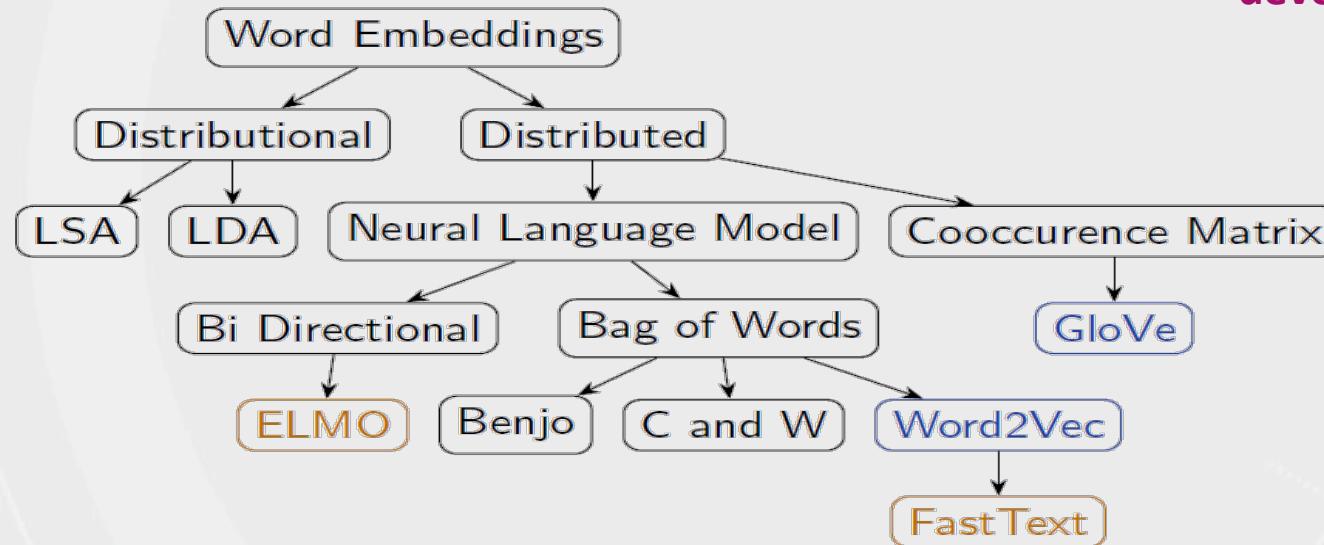
groom bride husband wife	Athens Greece Baghdad Iraq	short shorter small smaller
groom bride king queen	Athens Greece Bangkok Thailand	short shorter smart smarter
groom bride man woman	Athens Greece Beijing China	short shorter strong stronger
groom bride nephew niece	Athens Greece Berlin Germany	short shorter tall taller
groom bride policeman policewoman	Athens Greece Bern Switzerland	short shorter tight tighter
groom bride prince princess		short shorter tough tougher
		short shorter warm warmer
		short shorter weak weaker
		short shorter wide wider
		short shorter young younger
		short shorter bad worse



Word Embeddings: Techniques

Many ways of obtaining word vectors

New techniques are being developed every year!





Word Embeddings: Word2Vec(tor)

Word2Vec (Mikolov et al. 2013)

- For learning word representations using a shallow neural network
- No shared parameters across the network
- Fast training
 - Training over Wikipedia snapshot (about 1 billion words) with 50 dimensions takes only 4 hours
- Strong representations – previous state-of-the-art performance.
- Learned word embedding vectors can be represented as linear translations etc.

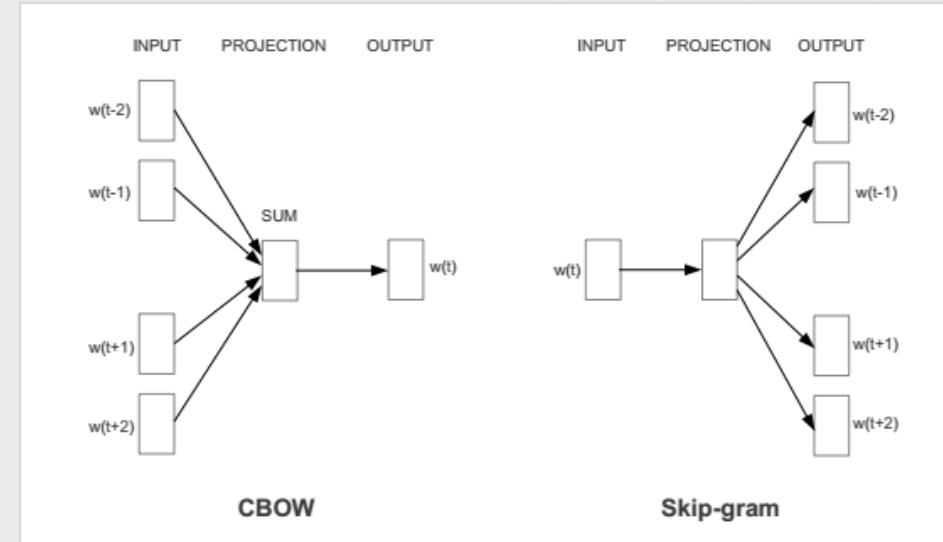
$$\begin{aligned}\text{vec(king)} - \text{vec(man)} + \text{vec(woman)} &\cong \text{vec(queen)} \\ \text{vec(Paris)} - \text{vec(France)} + \text{vec(China)} &\cong \text{Beijing}\end{aligned}$$



Word Embeddings: Word2Vec(itor) Continued.

Word2Vec Configurations

- Word2vec has two models
- Continuous Bag of Word (CBOW)
 - Predict a word $w(t)$ given its context
- Skip-gram models
 - Predict context words given a word $w(t)$

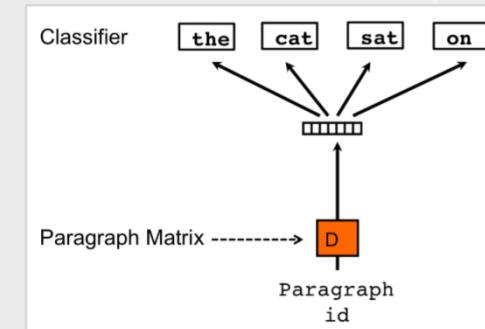
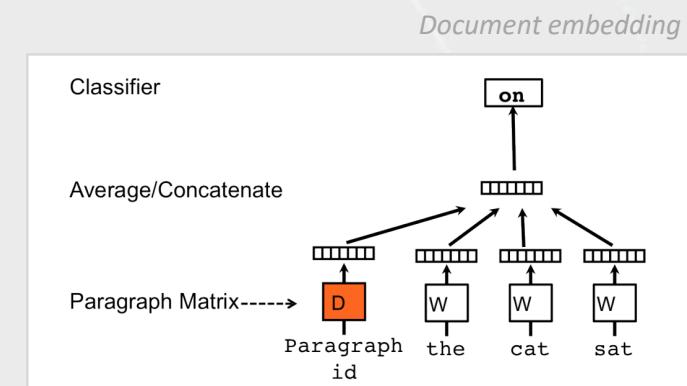




Other Language Representations

Different Embeddings Forms

- Sentences - Sent2Vec
- Paragraphs - Doc2Vec
- Document topics - Top2Vec
- Multi-Modal - Data2Vec





Fundamentals of Natural Language Processing

Agenda (Hands-on Session) – Word Representations

Learning word representations from scratch

- General domain and domain-specific word embeddings
- Word similarity
- Word vector clustering and visualisation

Supplementary Content

- Interactive exploration of word2vec models: <https://ronxin.github.io/wevi/>

12.3 Supervised Learning

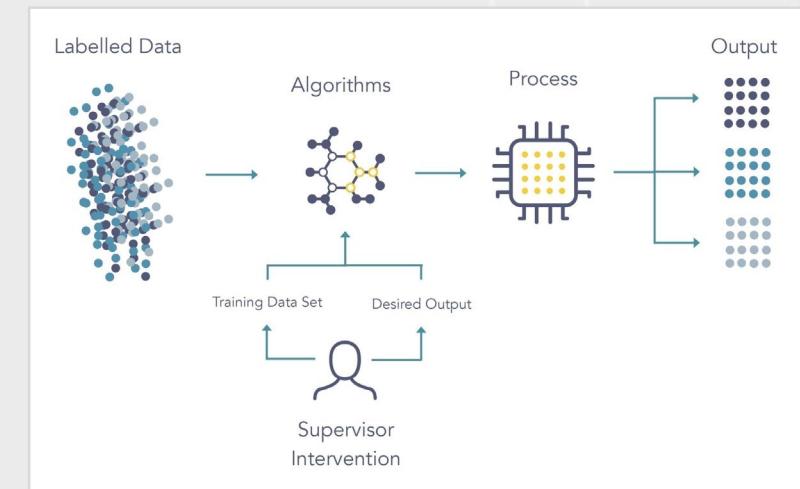
Learning from example



Supervised Learning

Fundamentals of Supervised Learning

- **Aim:** Learn a complex function that maps X to y when provided a set of labelled x, y pairs as example
- x, y pairs are typically acquired through human annotation using a pre-defined schema/model
- Can vary in difficulty and complexity





Supervised Learning

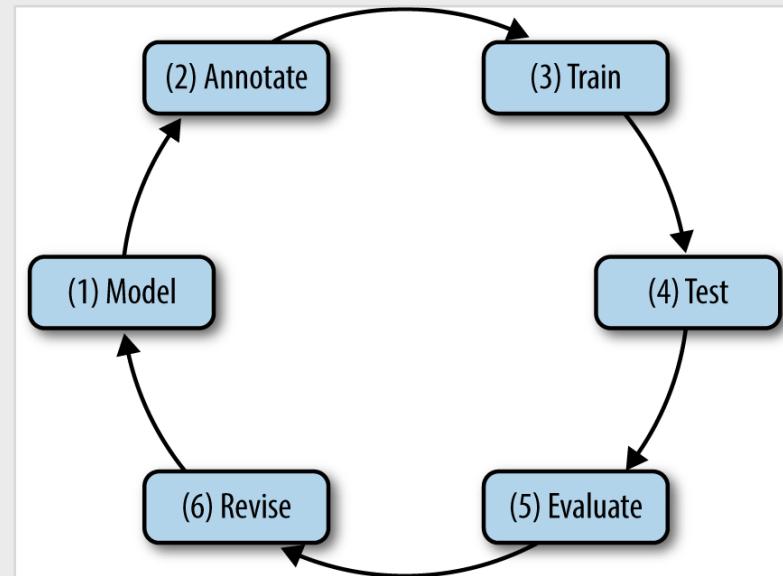
Agenda (Hands-on Session) – Exploring Supervised NLP Tasks using HuggingFace 

- Sequence Classification
- Question Answering
- Language Modelling
- Text Generation
- Named Entity Recognition
- Summarization
- Translation
- [Colab Notebook](#)



Fundamentals of Natural Language Annotation

The Annotation Development Cycle





Annotation Example: Accident Report (NER)

1 - Specifying a Model for Annotation

- What are the goals of the annotation task?
- What are the characteristics of the dataset?
Is it representative of the goals of the task?
Is it balanced? What is the quality?
- Does the task require subject matter experts or linguists?
- Does a complete or partial model already exist?

walking in car park from
office to workshop and
rolled ankle

injury

activity

body part

Initial Model



Annotation Example: Accident Report (NER)

2 - Performing human annotation

- What tool will be used to perform annotation?
- How many annotators are required to arrive at a *gold standard*?
- What level of resources are committed to the annotation task?
- How much data will be sufficient for training and testing?

injury

activity

body part

activity

walking in car park from
office to workshop and

rolled ankle

injury

body part



Annotation Example: Accident Report (NER)

3 & 4 - Training and Testing

- What model architecture is best suited for the chosen task?
- How will the texts be represented?
- How will the model architecture be optimised?
- What is the acceptable level of performance and expected bayes error rate?

injury

activity

body part

activity

walking in car park from

office to workshop and

rolled ankle

injury body part



Annotation Example: Accident Report (NER)

5 – Evaluation of Model Performance

- Does the model meet the acceptable level of performance?
- What does the model do well at, what does the model struggle with?

injury

activity

body part

activity

walking in car park from
office to workshop and

rolled ankle

injury

body part



Annotation Example: Accident Report (NER)

6 – Revising the Model

- Does the model meet the goals of the task? If not, what modifications need to be made?
- How are the annotators performing? Is the task too challenging?

walking in car park from
office to workshop and
rolled ankle

injury

activity

body part

location

New Model



Annotation Example: Accident Report (NER)

activity

location

location

location

walking in car park from office to workshop

and rolled ankle

injury

body part



Annotation Example: Accident Report (ET)

walking in car park from office to workshop

activity
activity/walking

location
location/car park

location
location/office

location
location/workshop

and rolled ankle

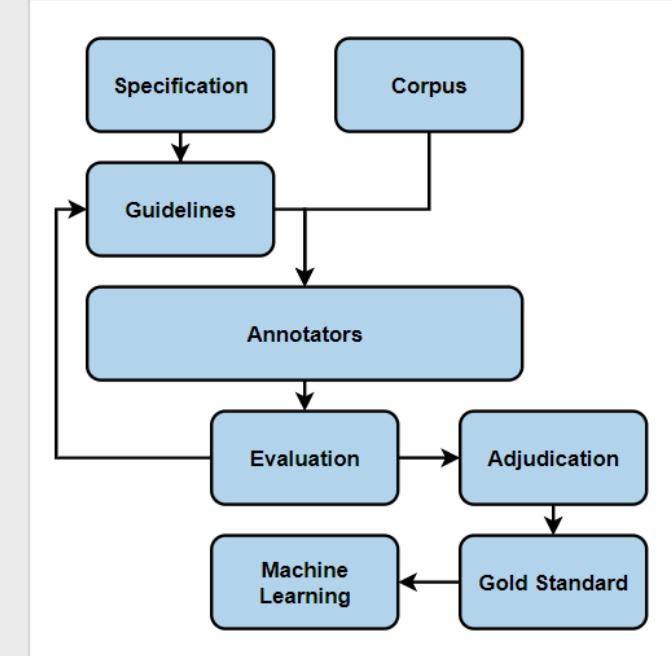
injury
injury/sprain body part
 body part/leg
 body part/leg/ankle



Fundamentals of Natural Language Annotation

After Development - The Annotation Cycle

- **Specification**: model of phenomena for a specific task
- **Guidelines**: helps annotators reliably label or tag the corpus
- **Annotators**: recruited or crowd sourced humans
- **Evaluation**: inter-annotator agreement
- **Adjudication**: creating a gold-standard annotated corpus
- **Machine Learning**: using annotations for supervised learning



12.4 Unsupervised Learning

Learning without example



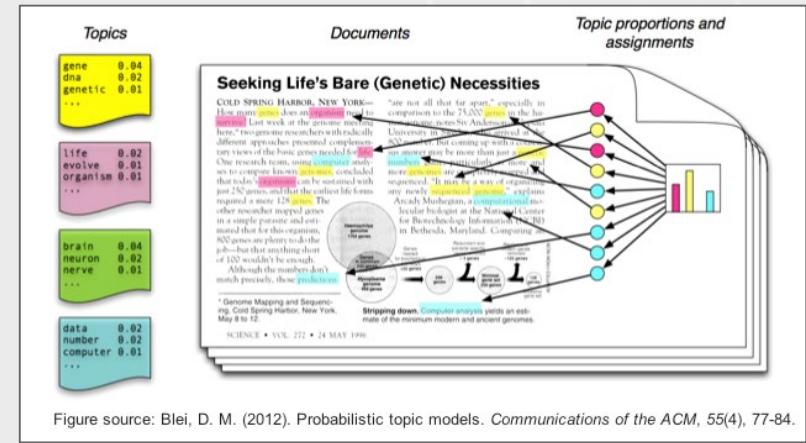
Unsupervised Learning: Topic Modelling

Suppose you have the following set of sentences:

- I eat **fish** and **vegetables**.
- *Fish* are *pets*.
- My *kitten* eats **fish**.

Latent Dirichlet allocation (LDA) is an unsupervised topic modelling algorithm that automatically discovers topics that these documents contain.

- **bold** words under the **Topic F**, which we might label as “**food**”
- *italics* words might be classified under a separate *Topic P*, which we might label as “*pets*”





Unsupervised Learning: Topic Modelling

LDA defines each topic as a bag-of-words (bow), and you have to label the topics as you deem fit.

1. We can infer the content spread of each sentence by a word count
 - **Sentence 1 (I eat fish and vegetables):** 100% Topic F
 - **Sentence 2 (Fish are pets):** 100% Topic P
 - **Sentence 3 (My kitten eats fish):** 33% Topic P and 67% Topic F
2. We can derive the proportions that each word constitutes in given topics. For example, Topic F might comprise words in the following proportions: 40% eat, 40% fish, 20% vegetables, ...



Unsupervised Learning: Topic Modelling

Three Steps

LDA defines each topic as a bag-of-words (bow), and you have to label the topics as you deem fit.

1. Specify how many latent topics are there
2. For each word, assign a temporary topic, according to a Dirichlet distribution, e.g. kitten -> F
 - If a word appears twice, each word may be assigned to different topics.
 - Function words (e.g. “the”, “and”, “my”) are removed and not assigned to any topics.
3. Topic assignment is updated based on two criteria:
 - How prevalent is that word across topics?
 - How prevalent are topics in the document?



Unsupervised Learning: Topic Modelling

Repeated refinement

- The process of checking topic assignment is repeated for each word in every document, cycling through the entire collection of documents multiple times.
- This iterative updating is the key feature of LDA that generates a final solution with coherent topics.

	Document X		Document Y
	Fish		Fish
	Fish		Fish
	Eat		Milk
	Eat		Kitten
	Vegetables		Kitten

	Document X		Document Y
F	Fish	?	Fish
F	Fish	F	Fish
F	Eat	F	Milk
F	Eat	P	Kitten
F	Vegetables	P	Kitten

	Document X		Document Y
F	Fish	?	Fish
F	Fish	F	Fish
F	Eat	F	Milk
F	Eat	P	Kitten
F	Vegetables	P	Kitten

	Document X		Document Y
F	Fish	?	Fish
F	Fish	F	Fish
F	Eat	F	Milk
F	Eat	P	Kitten
F	Vegetables	P	Kitten



Unsupervised Learning

Agenda (Hands-on Session) – Topic Modelling using Latent Dirichlet Allocation (LDA)

- Corpus pre-processing including lemmatization
- Training LDA model from scratch
- Interactive visualisation of LDA model
- Exploration of LDA features



COREHUB.COM.AU/SKILLS

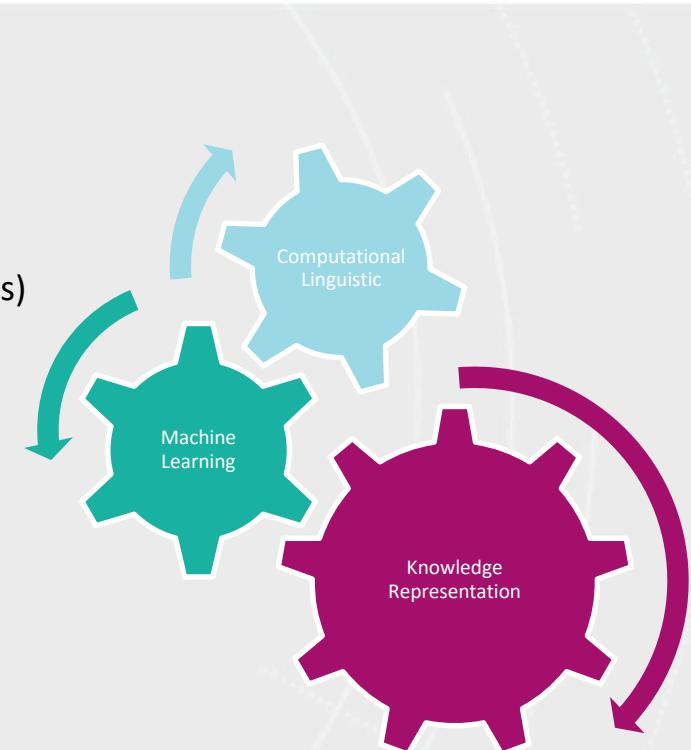


Additional Content



NLP Pipelines – Ontology Learning from Text

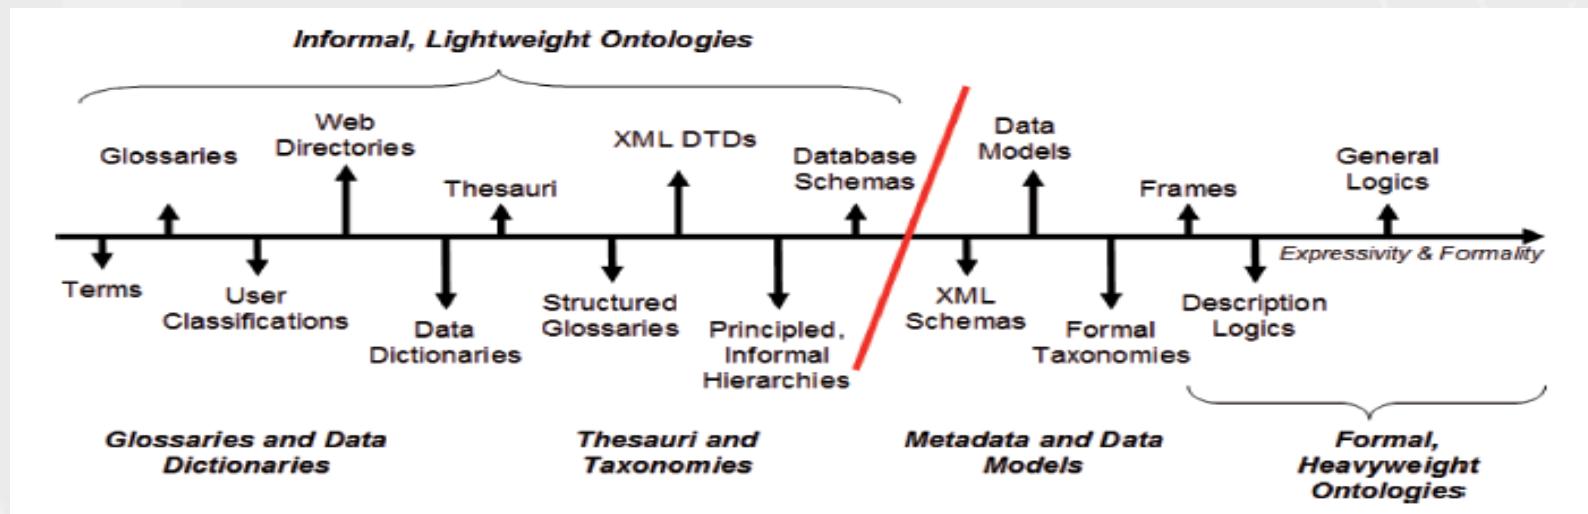
- **Graphs in knowledge representation**
 - Two different sources
 - Knowledge experts (logic based)
 - Knowledge discovery (from natural language texts)
- Ontology learning
 - Entity Extraction
 - Symbolic
 - Unsupervised learning
 - Supervised learning
 - Subsymbolic
 - Deep learning
 - Entity Relation Extraction





Lightweight Ontologies

Overview – What's Light Weight Ontology



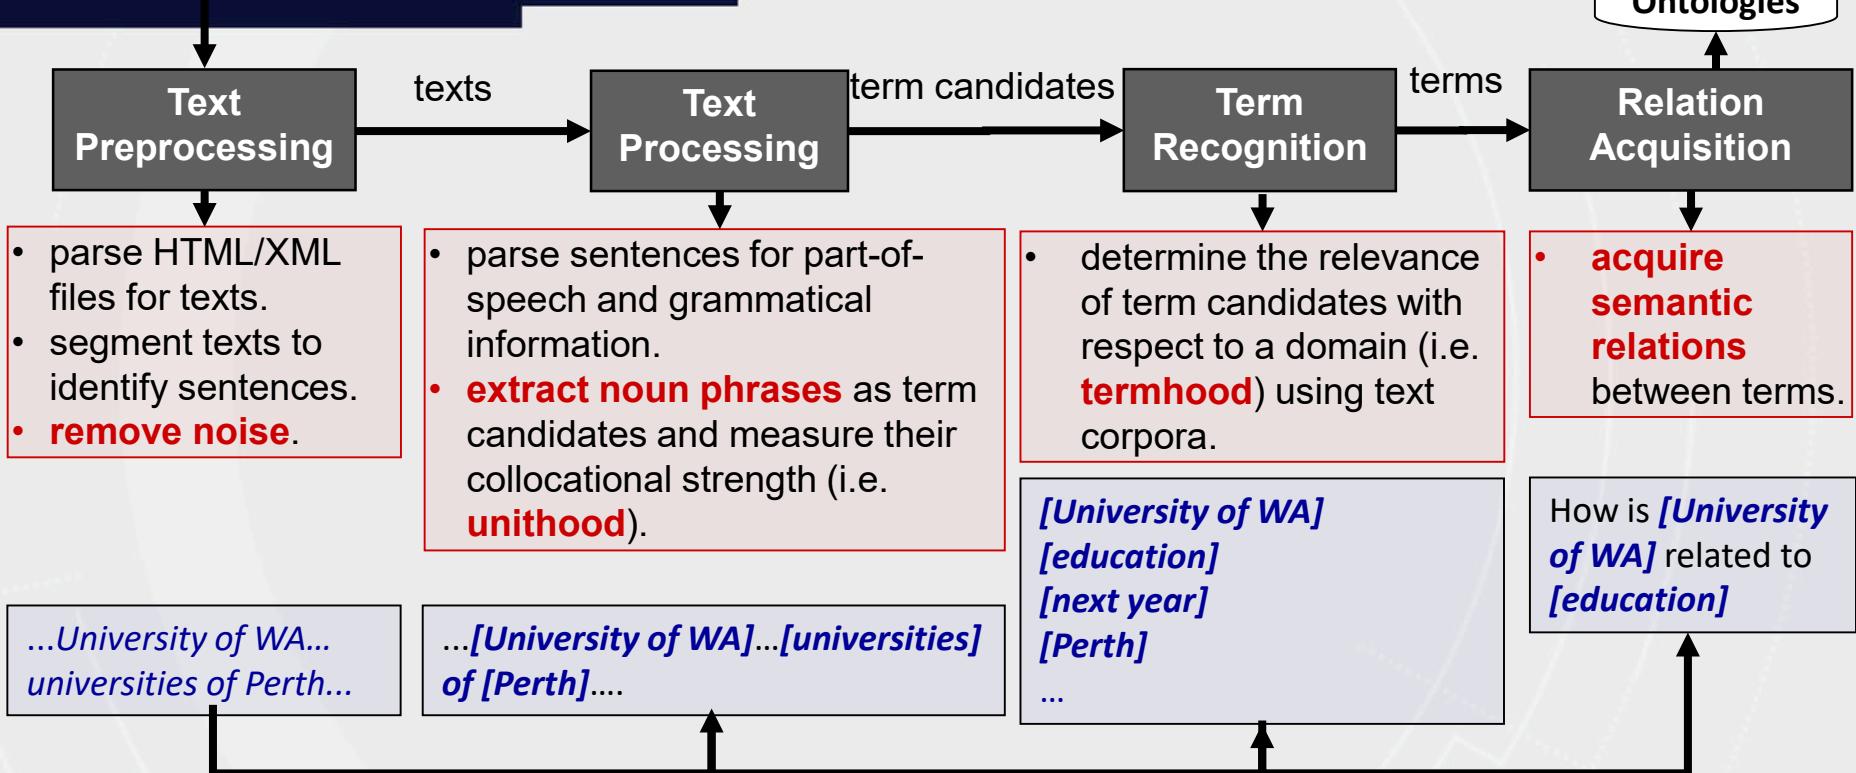
Wilson Wong, Wei Liu, and Mohammed Bennamoun. (2012) *Ontology Learning from Text: A look back and into the future*, ACM COMPUTING SURVEYS, 44, 4, Article 20, pp. 20:1-20:36

Associate Prof. Wei Liu, UWA





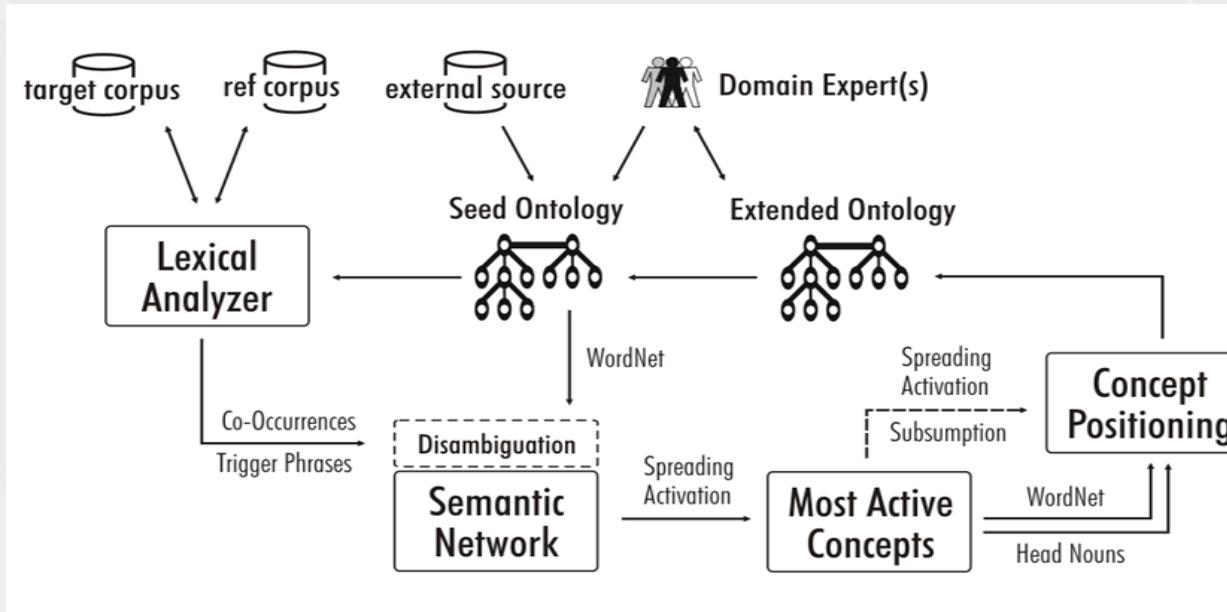
text > sets of corpora are prepared for ontology construction. Web Crawling is performed to gather web pages from general sources such as Reuters, Discovery and CNet to create the Contrastive Corpus. Readily available non-domain specific





Ontology Learning from Text – System Overview

A semi-automatic ontology learning system



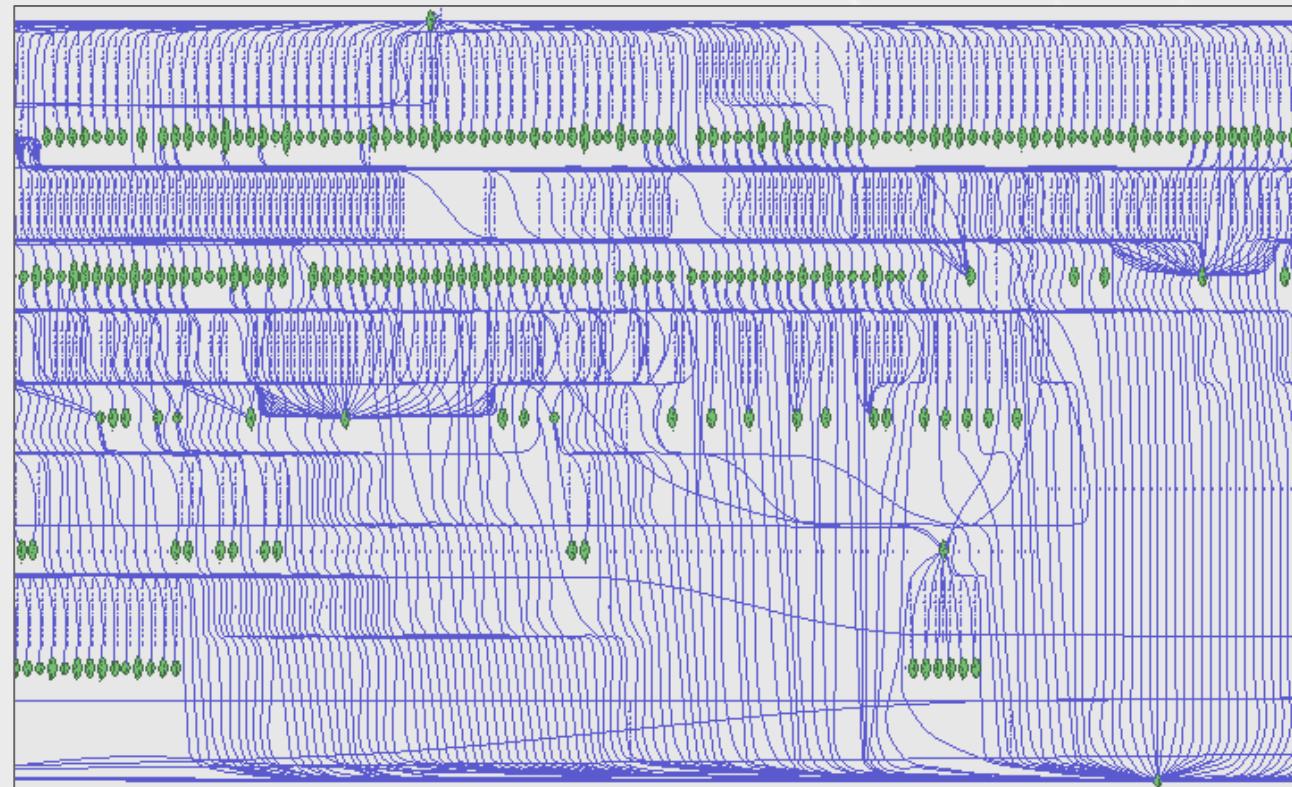
Wei Liu, Arno Scharl & Albert Weischselbraun, 2005

Associate Prof. Wei Liu, UWA



Ontology Learning from Text – Term Co-occurrence

Semantic Network
Visualisation





Ontology Learning from Text – Term Significance

Results of climate change data analysis

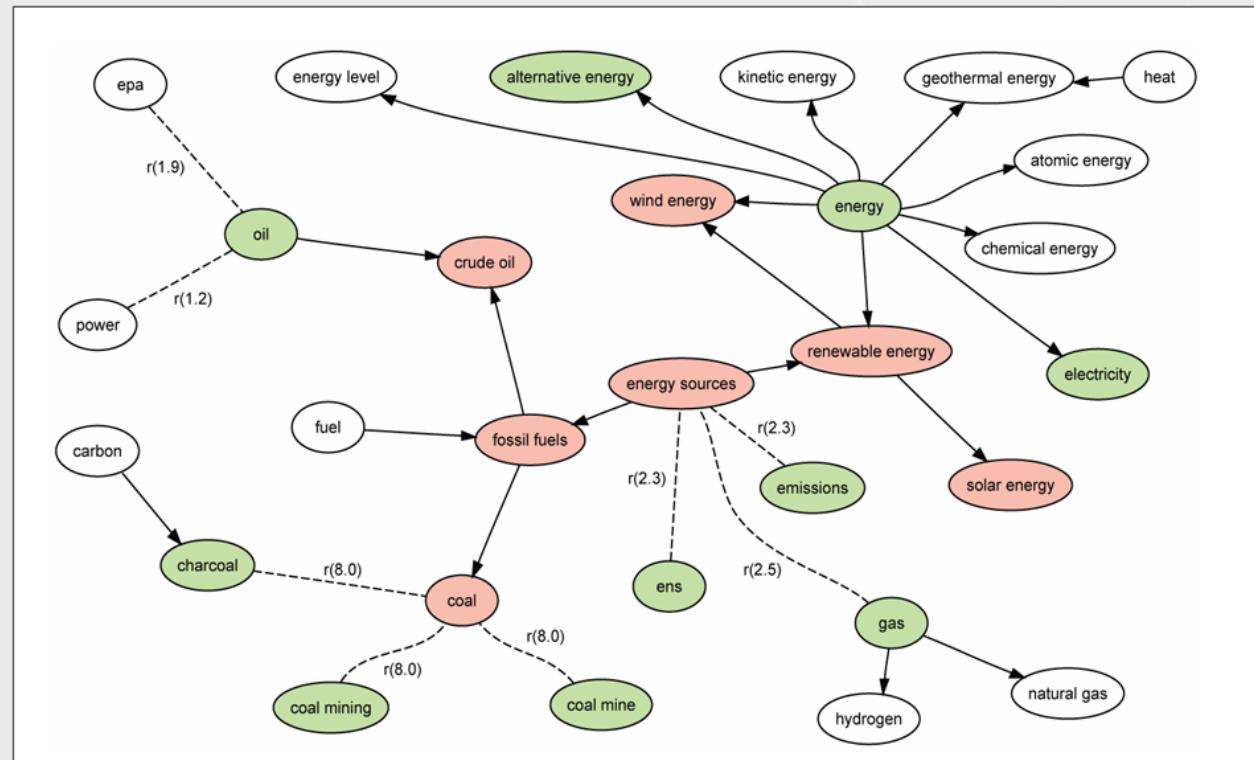
ENERGY		
DOCUMENT LEVEL	SIG	
gas	128,714	
power	67,176	
natural gas	62,396	
natural	46,337	
electricity	43,735	
fuel	25,761	
oil	25,168	
allegheny	24,183	
renewable	23,815	
utility	23,554	

ENERGY		
SENTENCE LEVEL	SIG	
renewable energy	204,328	
renewable	194,479	
allegheny energy	185,018	
allegheny	154,046	
energy efficiency	147,639	
duke energy	143,573	
energys	137,306	
sempra	119,689	
sempra energy	105,757	
centerpoint	101,068	



Ontology Learning from Text – Building Lightweight Ontologies

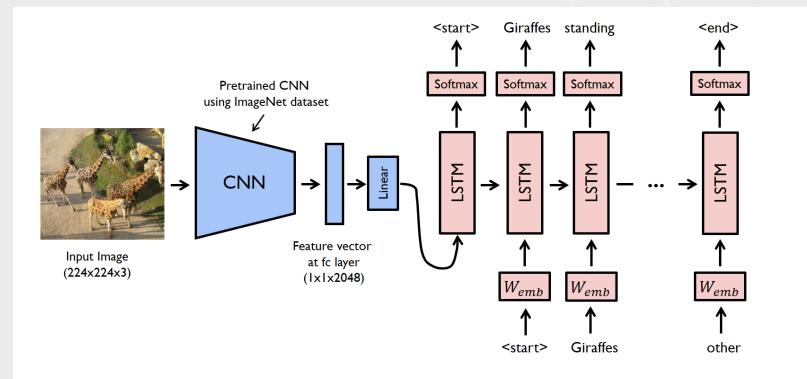
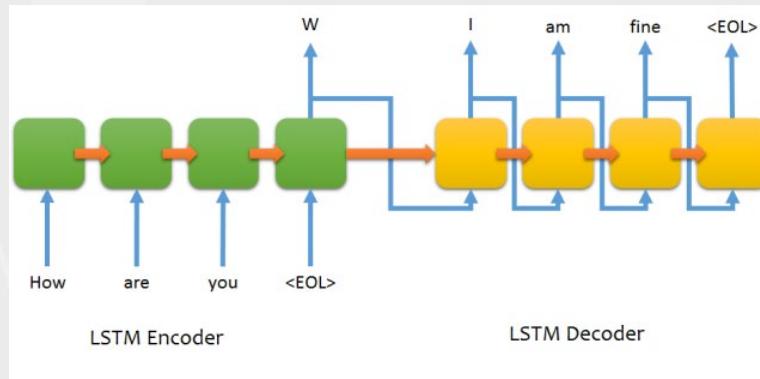
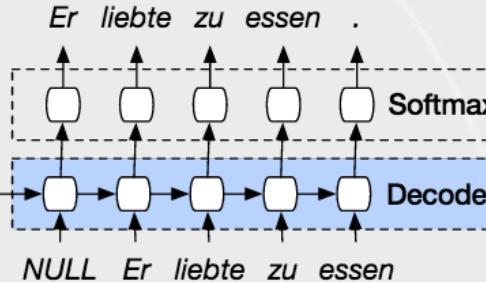
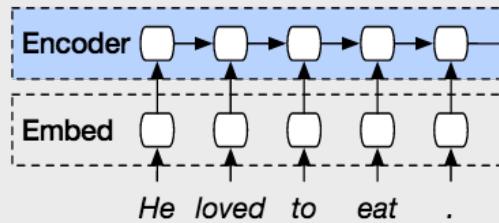
Spreading Activation



12.4 Neural Language Processing



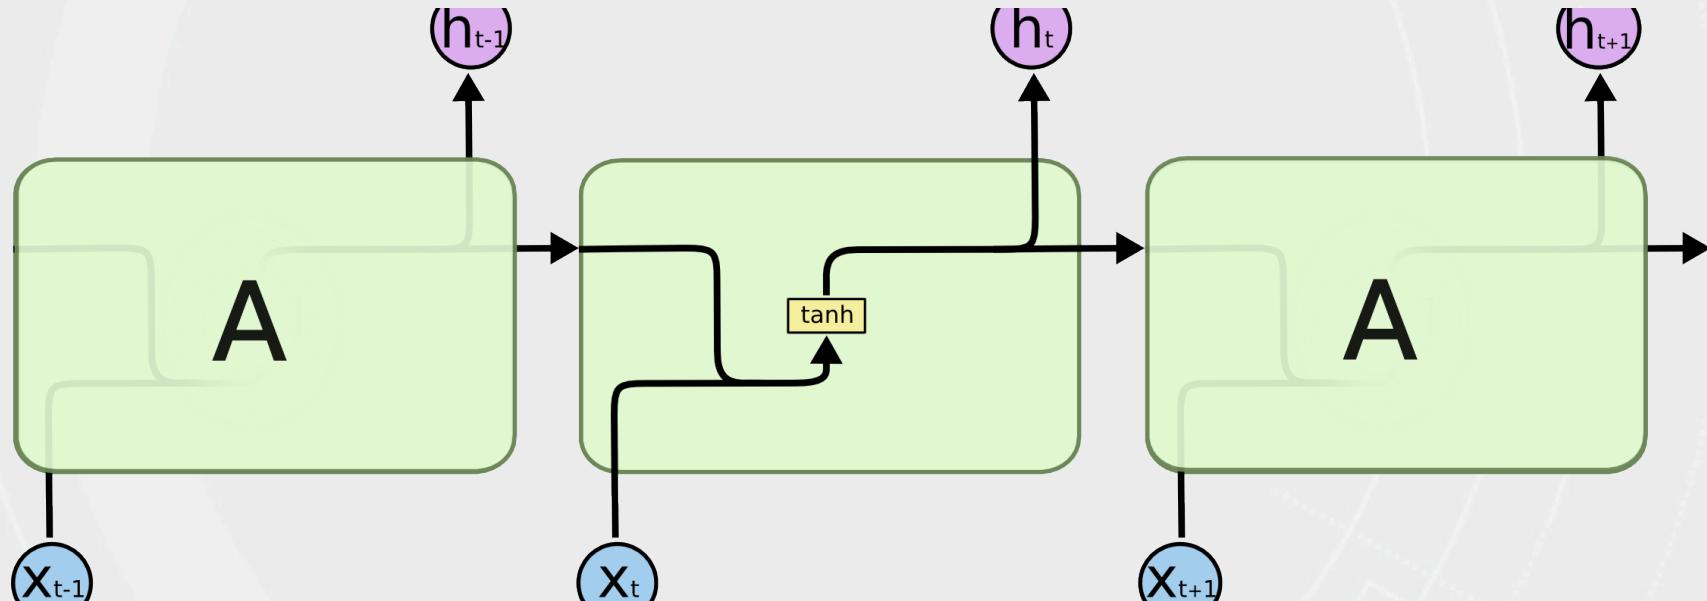
Sequence to Sequence Learning (Encode then Decode)





LSTM

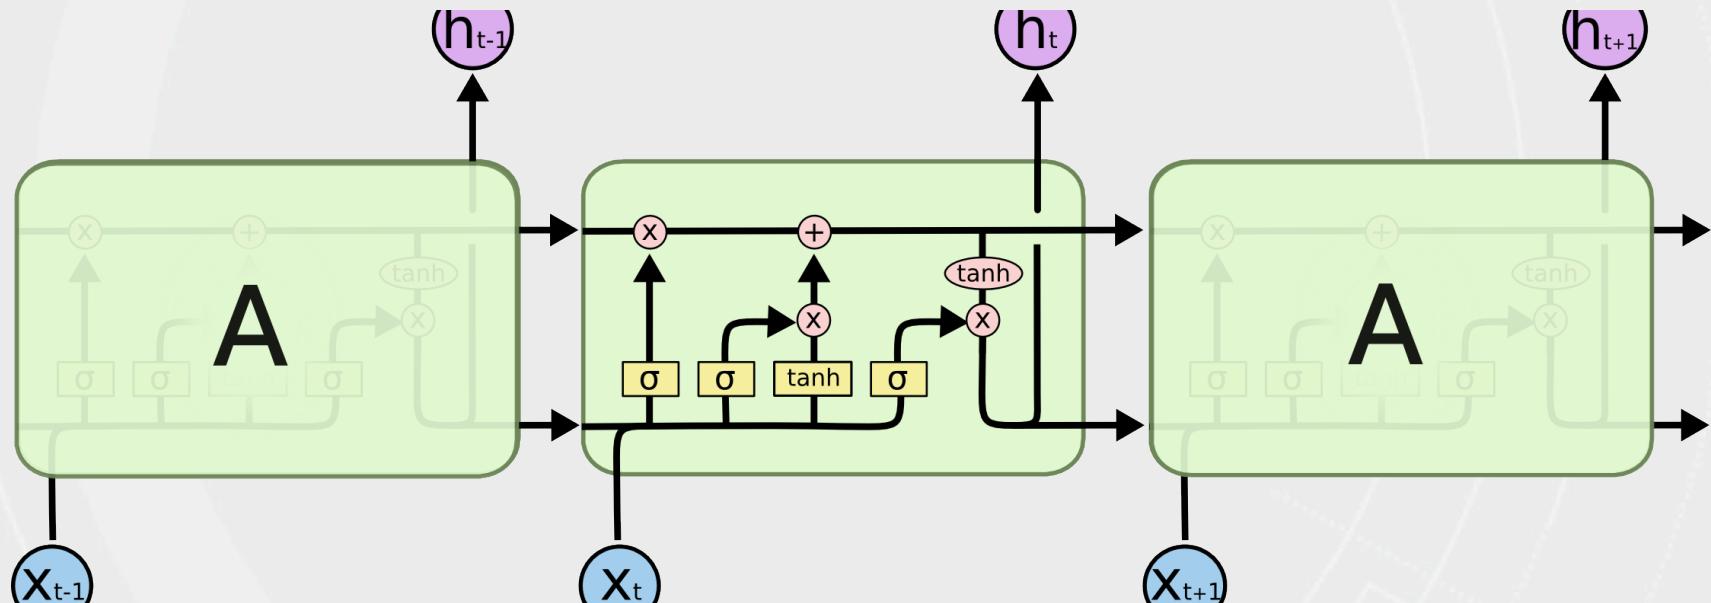
Contains four interacting layers





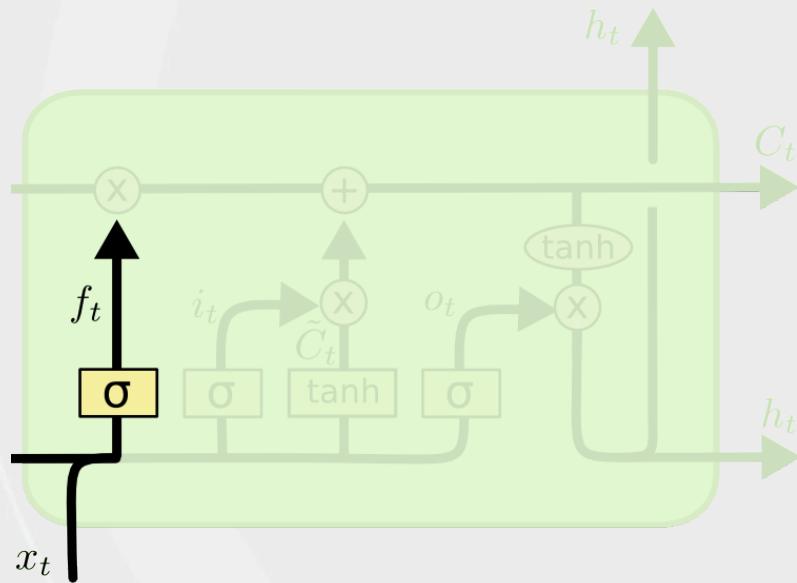
LSTM

Inside a cell





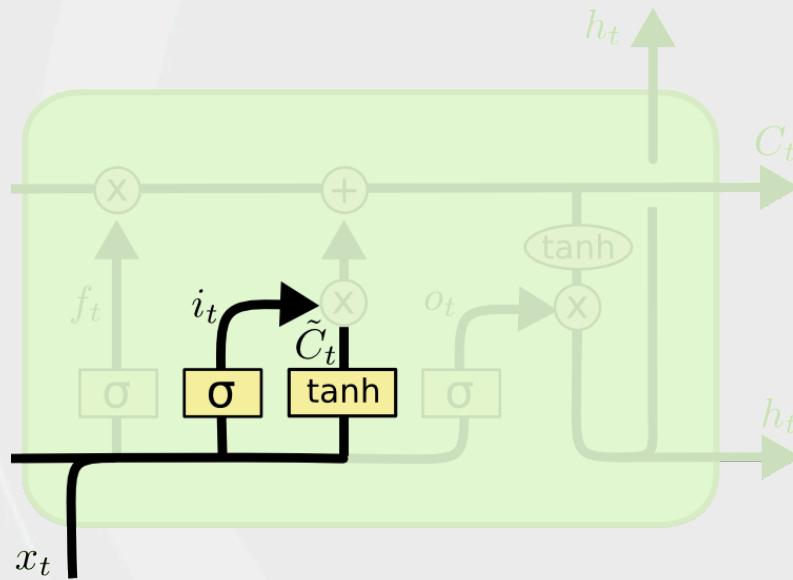
LSTM – Forget gate



$$f_t = \sigma (W_f \cdot [h_{t-1}, x_t] + b_f)$$



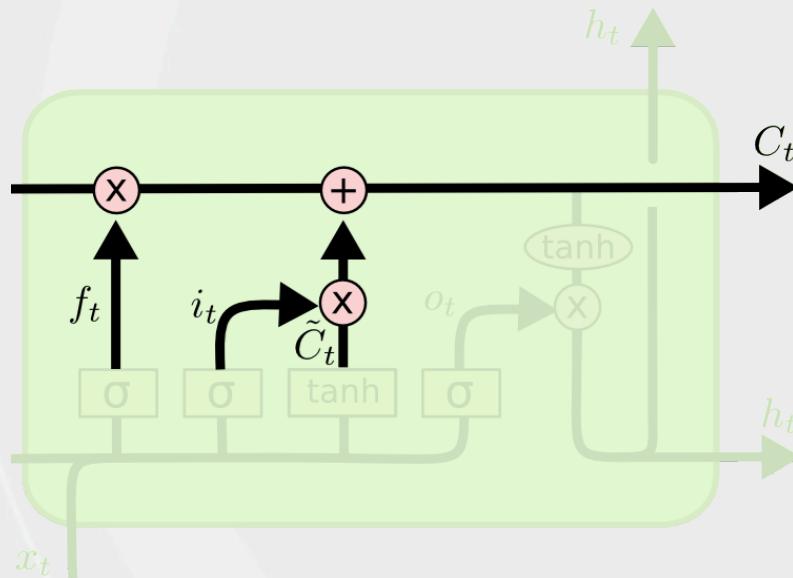
LSTM – Input Gate Layer



$$i_t = \sigma (W_i \cdot [h_{t-1}, x_t] + b_i)$$
$$\tilde{C}_t = \tanh(W_C \cdot [h_{t-1}, x_t] + b_C)$$



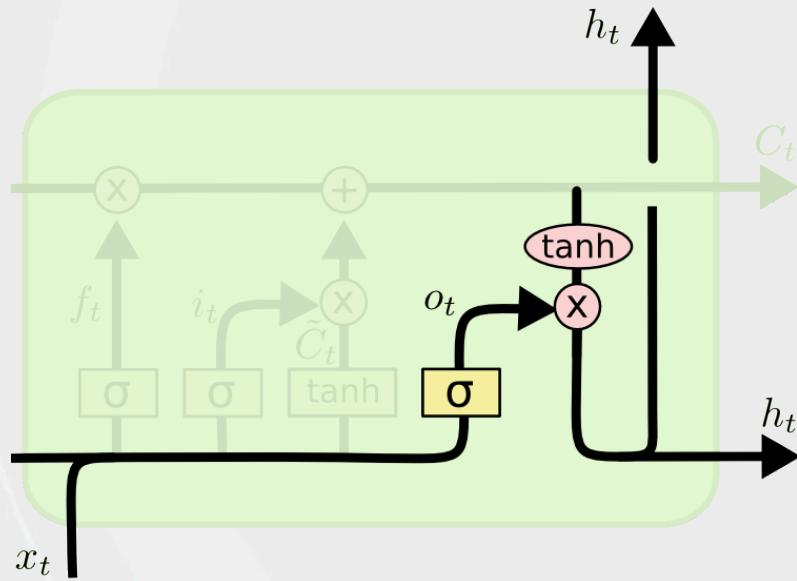
LSTM - Update



$$C_t = f_t * C_{t-1} + i_t * \tilde{C}_t$$



LSTM – Output Gate

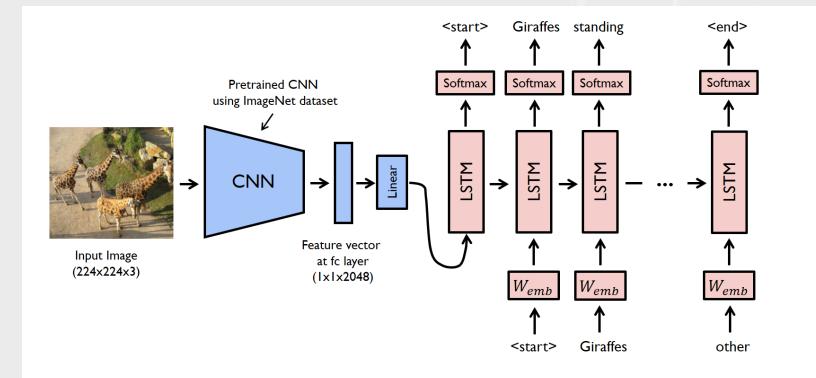
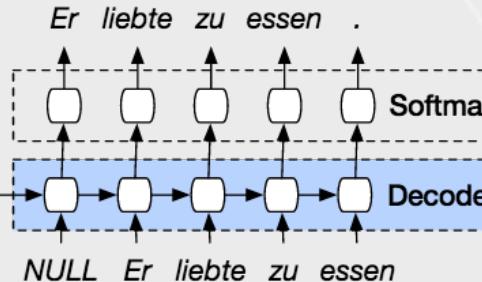
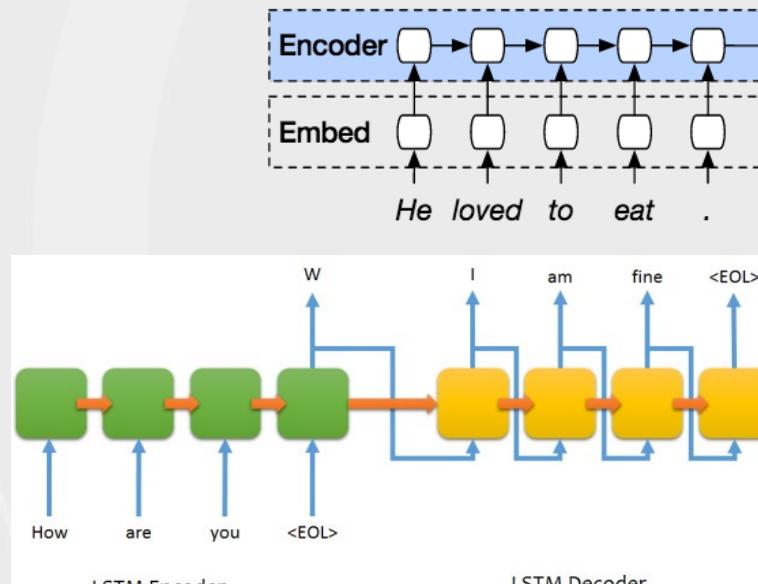


$$o_t = \sigma (W_o [h_{t-1}, x_t] + b_o)$$

$$h_t = o_t * \tanh (C_t)$$



Sequence to Sequence Learning (Encode then Decode)





Associate Prof. Wei Liu

Educator



Program Chair of Data Science, Department of Computer Science and Software Engineering, The University of Western Australia

CORE Skills Data Science Springboard Delivery Team

Wei Liu is a teaching and research academic at UWA in the Department of Computer Science and Software Engineering focusing on knowledge discovery from natural language text. Her research focuses on knowledge discovery from natural language text, deep learning methods for knowledge graph construction and analysis, as well as sequential data mining and forecasting in traffic and water consumption domain.

Wei received her PhD from the University of Newcastle, Australia in 2003. She has published in highly reputable journals such as ACM Computer Surveys, Journal of Data Mining and Knowledge Discovery, Knowledge and Information Systems and presented at key international events including the International Conference on Data Engineering (ICDE) and ACM International Conference on Information and Knowledge Management (CIKM).

Wei has won three Australian Research Council Grants and managed several industry grants. Her current industry-related research projects include knowledge graph refinement for geological survey reports, incident log analysis and visualisation, short-term traffic predication and cognitive computing for asset management.

Program Delivery Day 12 (Special Data Types - Natural Language Processing & Text Mining)



Tyler Bikaun

Educator



PhD Candidate (Technical Language Processing), Department of Computer Science and Software Engineering, The University of Western Australia & Mechanical Engineer

CORE Skills Data Science Springboard Delivery Team

Tyler Bikaun is a Mechanical Engineer turned PhD candidate at UWA and a scholar with the Mineral Research Institute of Western Australia and ARC Training Centre for Transforming Maintenance through Data Science. His research focuses on the intersection of computer science and industrial engineering, in particular, applying computational techniques to technical texts in the context of industrial maintenance.

Prior to commencing his PhD, Tyler worked in engineering consulting to the resources and utilities sectors in Western Australia. Here he was involved in asset performance optimisation of rotating equipment and advanced analytics for asset health monitoring applications.

Program Delivery Day 12 (Special Data Types - Natural Language Processing & Text Mining)